



Module 02: Persistence API

CS544: Enterprise Architecture



Object Persistence

- So far we've only discussed the basics of mapping classes to tables
- In this module we will discuss Hibernate's object management
- We will look at the different states of the entity lifecycle, and the methods that are used to move objects between them
- We will also provide insight into the differences between some of Hibernate's seemingly duplicate methods

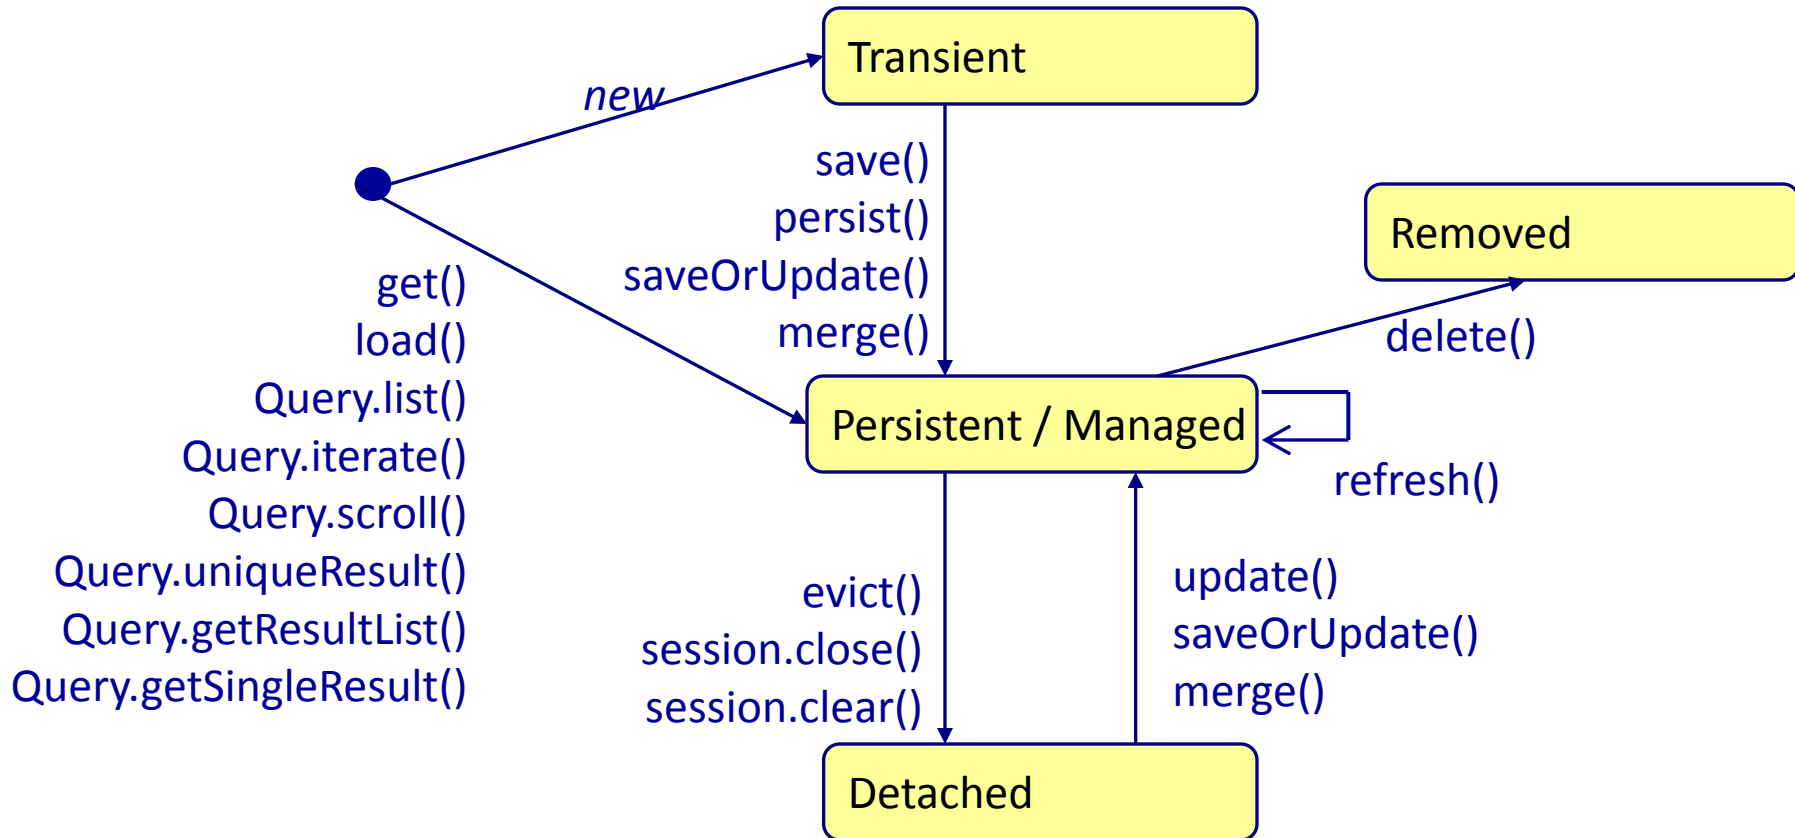


Persistence API:

ENTITY OBJECT LIFECYCLE



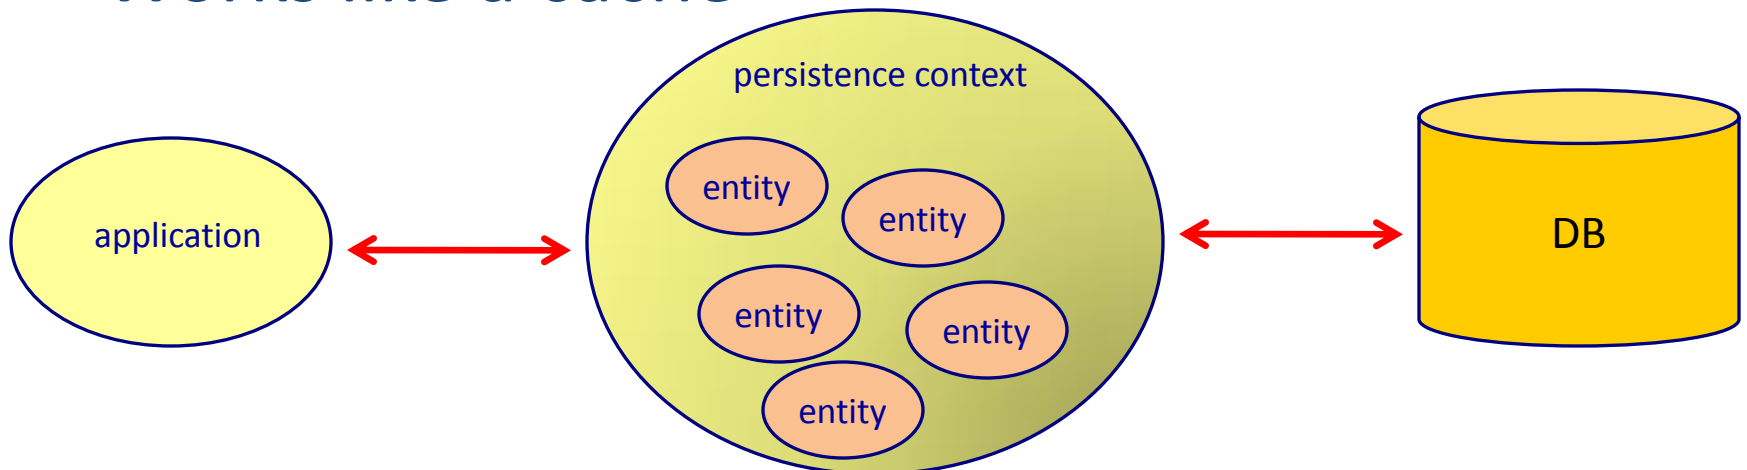
Lifecycle of an entity





Persistence context

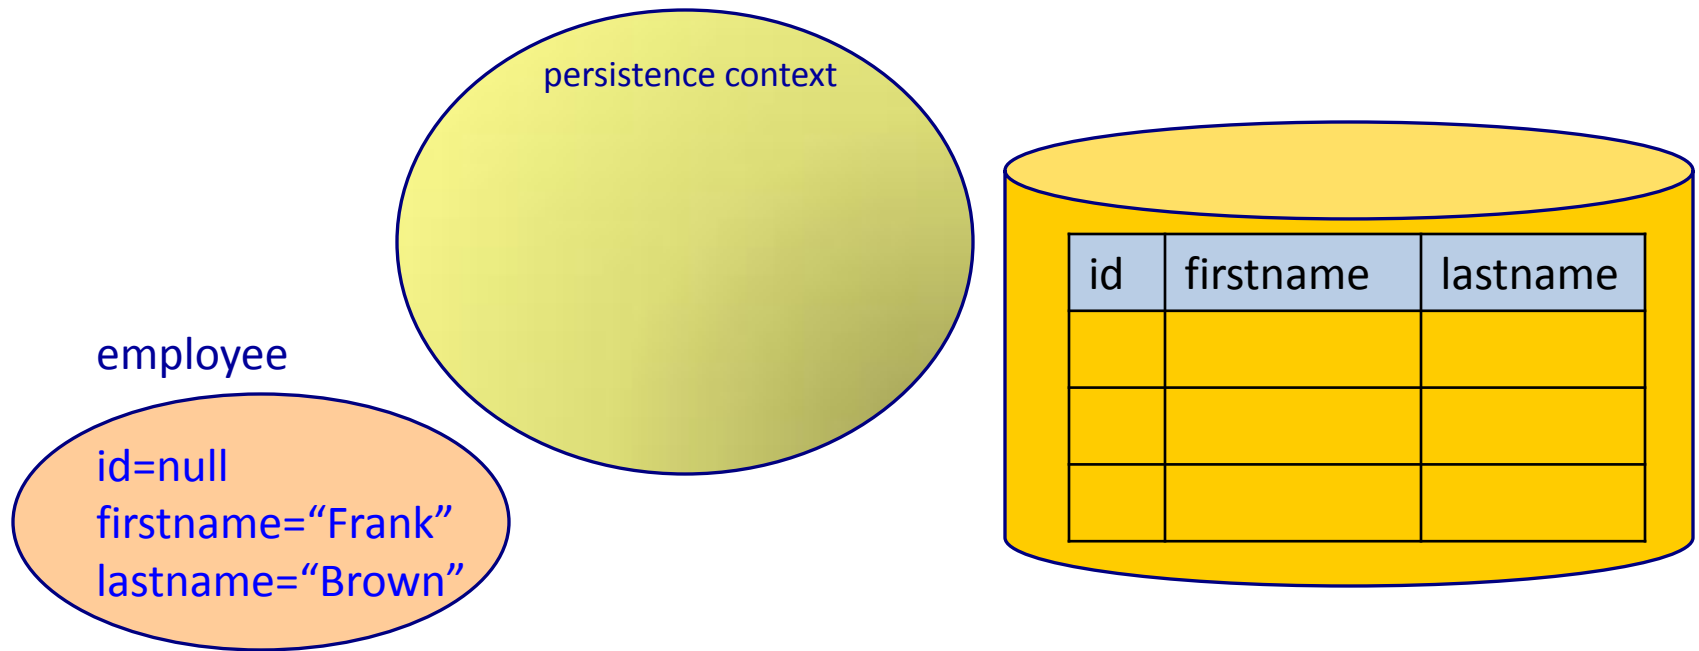
- Manages the entities
- Guarantees that managed entities will be saved in the database
- Tracks changes until they are pushed to the database
- Works like a cache





Transient entity


- A transient entity has no database identity





Managed entity

- A managed entity is managed by the persistence context and has a database identity



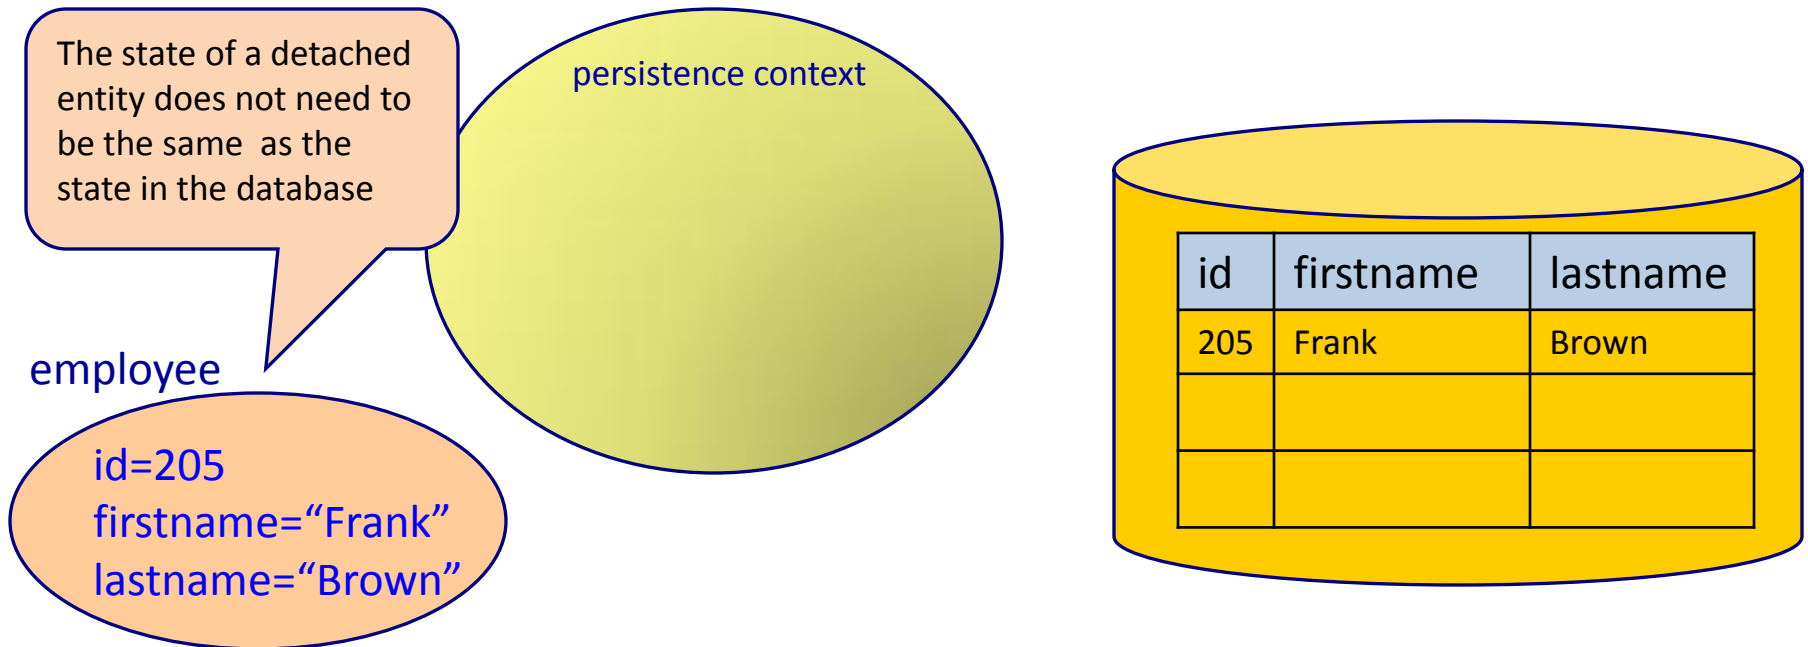
The diagram shows a yellow cylinder representing a database table. Inside the cylinder is a table with three columns: "id", "firstname", and "lastname". The first row contains the values "205", "Frank", and "Brown". The second and third rows are empty.

id	firstname	lastname
205	Frank	Brown



Detached entity

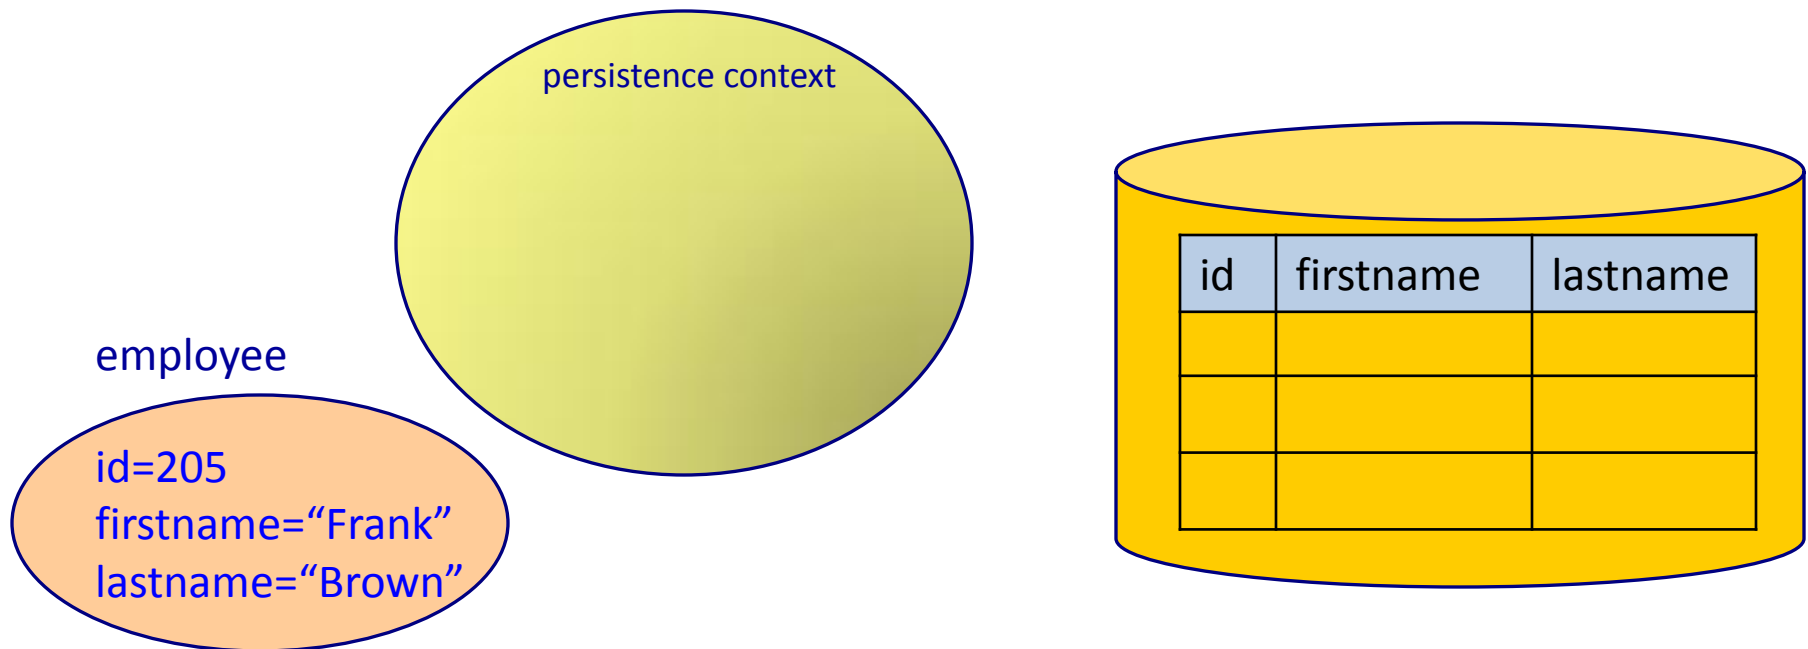
- A detached entity has a database identity, but is not managed by the current persistence context





Removed entity

- With a removed entity is the corresponding data removed from the database.





Session - Persistence Context

Type	Methods	Description
CREATE	<code>session.persist()</code> <code>session.save()</code>	Uses SQL INSERT statements to create rows in the database
RETRIEVE	<code>session.load()</code> <code>session.get()</code>	Uses SQL SELECT statements to retrieve rows from the database
UPDATE	*implicit update* <code>session.update()</code>	Uses SQL UPDATE statements to update rows in the database
CREATE / UPDATE	<code>session.saveOrUpdate()</code> <code>session.merge()</code>	Uses either INSERT or UPDATE depending on the state of the object
DELETE	<code>session.delete()</code>	Uses SQL DELETE statements to delete rows from the database
Cache related functions	<code>session.refresh()</code> <code>session.flush()</code>	Explicitly gets changes from the database, explicitly pushes changes to the database



Persistence API:

STORING ENTITY OBJECTS



About: save() and persist()

- Save() and persist() are very similar
 - Both will add the given entity to the persistence context, and INSERT it into the database

```
Person p1 = new Person("Frank Brown");  
long id1 = (Long)session.save(p1);
```

Saves the entity and
returns its generated id

```
Person p1 = new Person("Frank Brown");  
session.persist(p1);
```

Persists the entity but
does not return anything

- Save() will always insert **right away (eager)**
- Persist() will insert when it needs to **(lazy)**



Return value: save() / persist()

- Because **save()** executes right away it can guarantee to **return the generated id**
- **Persist()** on the other hand may wait until the end of the transaction to execute its INSERT
- **Persist()** therefore **returns void**

```
Person p1 = new Person("Frank Brown");  
long id1 = (Long)session.save(p1);
```

Saves the entity and
returns its generated id

```
Person p1 = new Person("Frank Brown");  
session.persist(p1);
```

Will persist the entity
but may not do so now,
can therefore not
guarantee a return id



Persistence API:

RETRIEVING ENTITY OBJECTS



About: get() and load()

- Get() and load() are very similar
 - Both retrieve an object by its primary key
 - `SELECT * FROM ... WHERE [primary key] = ...`
 - If an entity is already cached the database is not hit

```
Person p1 = (Person) session.get(Person.class, 1L);
```


```
Person p1 = (Person) session.load(Person.class, 1L);
```

- Get() retrieves the object's values **right away**
- Load() **provides a proxy** and does not retrieve the object's values until they are first needed




get()

1. `Product product = (Product)session.get(Product.class, new Long(productid));`




2. `product.getPrice()`



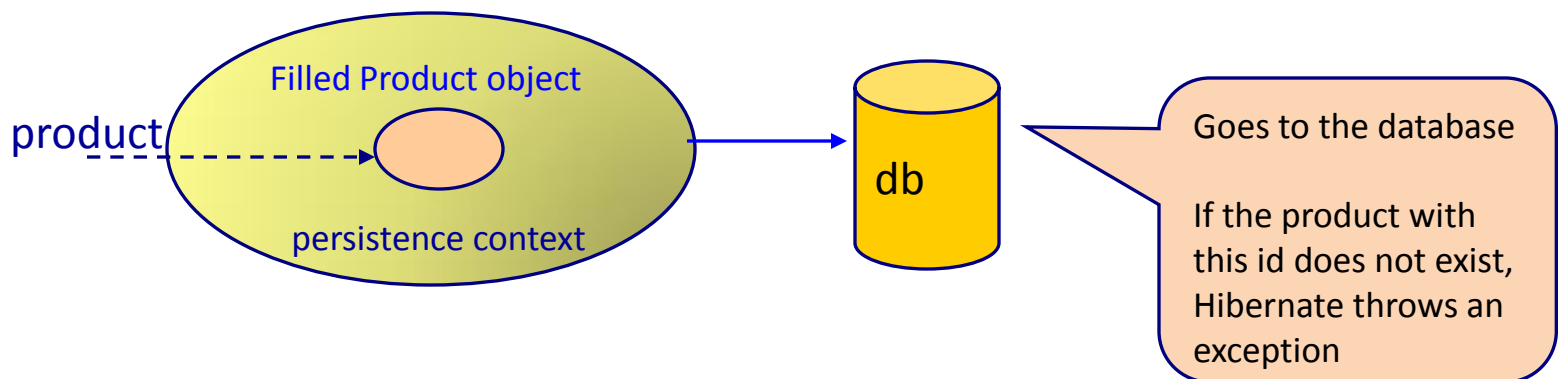


load()

1. `Product product = (Product)session.load(Product.class, new Long(productid));`



2. `product.getPrice()`





When to use load() instead of get()

```
Employee employee = (Employee) session.get(Employee.class, employeeid );  
session.delete(employee);
```

- 2 database hits
 - 1 SELECT statement for the get() method
 - 1 DELETE statement for the delete() method

The load() method creates an Employee proxy object in the persistence context, set its id, but does not hit the database.

```
Employee employee = (Employee) session.load(Employee.class, employeeid );  
session.delete(employee);
```

- 1 database hit
 - 1 DELETE statement for the delete() method



Persistence API:

UPDATING ENTITY OBJECTS



Implicit Update

- When a **managed entity is changed** inside a transaction, the changes are pushed to the database when the **transaction commits**

```
try {  
    session = sessionFactory.openSession();  
    tx = session.beginTransaction();  
    Person p3 = (Person) session.get(Person.class, 1L);  
    p3.setName("Implicitly Updated");  
    tx.commit();  
} catch (HibernateException e) {  
    tx.rollback();  
    e.printStackTrace();  
} finally {  
    if (session != null) {  
        session.close();  
    }  
}
```

Get () puts the object into the persistence context (becomes managed)

Object is changed

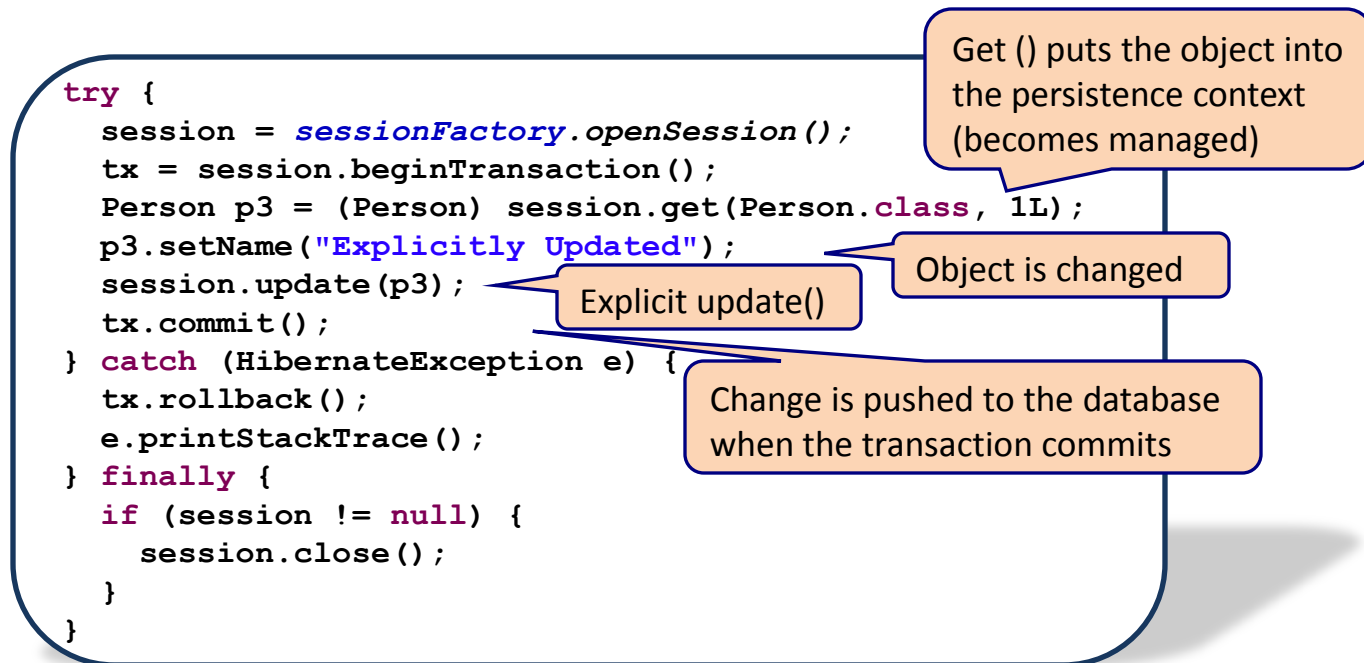
Change is pushed to the database when the transaction commits

- Hibernate notices that an object is 'dirty' and uses a SQL UPDATE to push it to the database



About: update()

- You can also explicitly call the update() method



- Like the implicit update, changes are not pushed to the database until the transaction is committed



Implicit vs. Explicit update

- In the last example the **explicit update** was completely redundant – **not needed**
- Why even have an `explicit update()` method?

```
session.update(detached_object);
```

Update() can be used to re-attach

- The real **power of the explicit update()** is in **re-attaching detached objects** that have been updated outside of the persistence context
- Once attached their changes will also be pushed to the database on `commit()`



Re-Attach Example

Session.close()
makes p1 detached

```
session = sessionFactory.openSession();  
Person p1 = new Person("Frank Brown");  
session.save(p1);  
session.close();
```

```
p1.setName("Detached Update");
```

Detached update

Update re-attaches

```
try {  
    session = sessionFactory.openSession();  
    tx = session.beginTransaction();  
    session.update(p1);  
    tx.commit();  
} catch (HibernateException e) {  
    tx.rollback();  
    e.printStackTrace();  
} finally {  
    if (session != null) {  
        session.close();  
    }  
}
```

Commit pushes updates of
attached object to the db

Load object to check if
update was successful

```
session = sessionFactory.openSession();  
Person p2 = (Person) session.get(Person.class, 1L);  
System.out.println("Name: " + p2.getName());  
session.close();
```

Person was updated
successfully

Name: Detached Update



The Problem with Re-Attaching

- Update() will throw a **NonUniqueObjectException**
 - when re-attaching an object with the same id as an already managed object

When update() receives a detached object for an entity it is already managing Hibernate will throw a NonUniqueObjectException!

employee

Id=105
firstname="Frank"
lastname="Update"

Detached object

persistence context
employee

id=105
firstname="Frank"
lastname="Brown"

Existing, already managed object

id	firstname	lastname
205	Frank	Brown

Solution: use merge()



Exception Example

```
session = sessionFactory.openSession();
Person p1 = new Person("Frank Brown");
long id = (Long)session.save(p1);
System.out.println("Frank Brown Id: " +id);
session.close();
```

Session.close()
makes p1 detached

Output shows id=1

```
try {
    session = sessionFactory.openSession();
    tx = session.beginTransaction();

    Person p2 = (Person)session.get(Person.class, 1L);
    session.update(p1);

    tx.commit();
} catch (HibernateException e) {
    tx.rollback();
    e.printStackTrace();
} finally {
    if (session != null) {
        session.close();
    }
}
```

Tries to re-attach
p1, also id=1
(Also Frank Brown)

New session, new
persistence context

Loads person with
id=1 (Frank Brown)
into the persistence
context

Output:

Hibernate: /* insert entity.Person */ insert into Person (id, name) values (null, ?)

Hibernate: call identity()

Frank Brown Id: 1

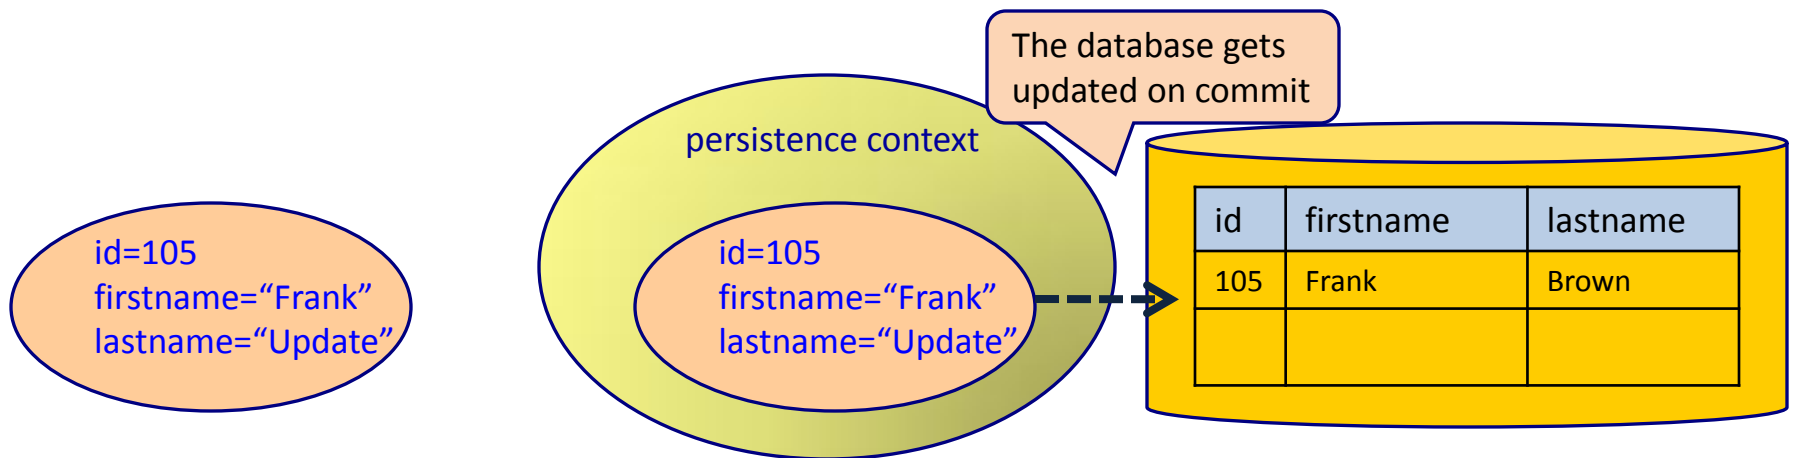
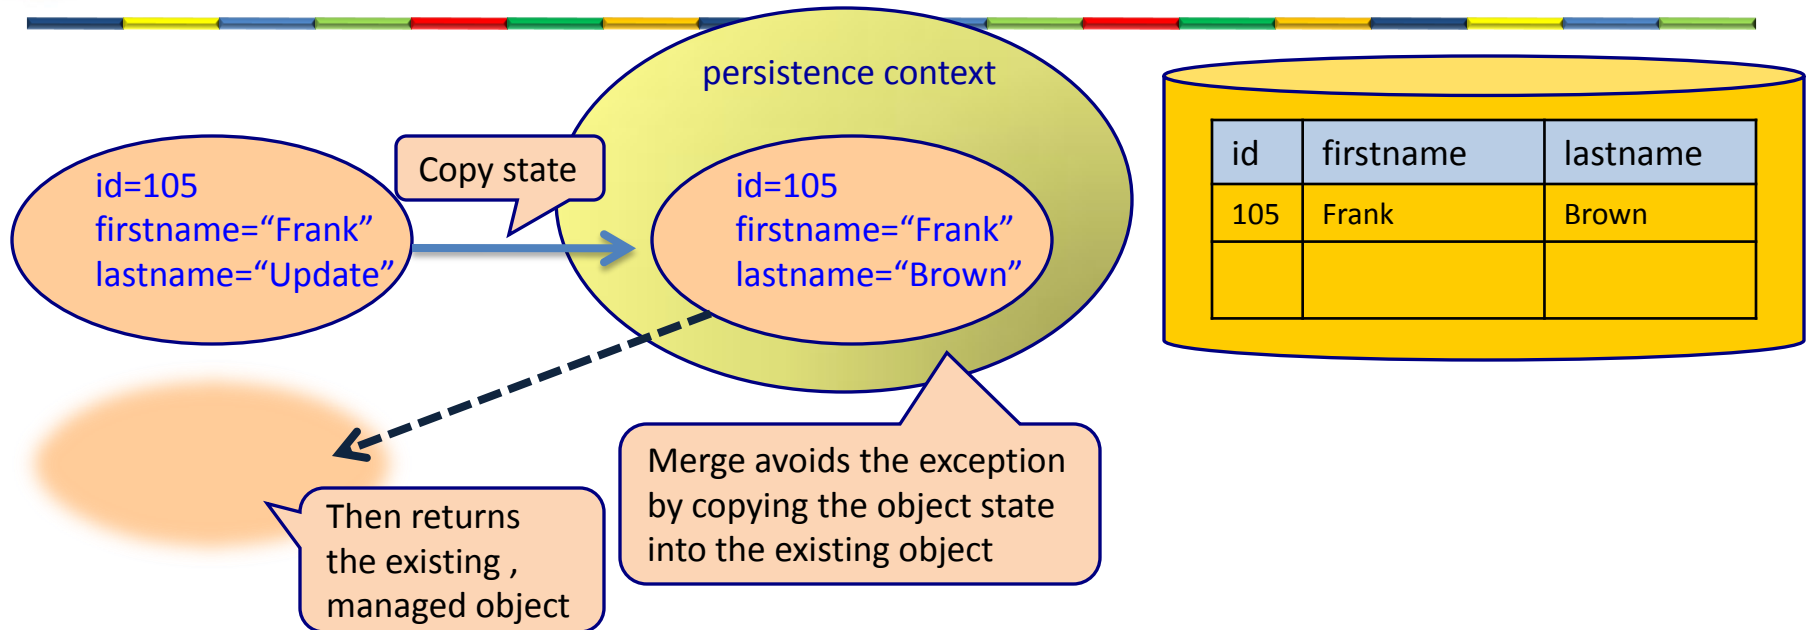
Hibernate: /* load entity.Person */ select person0_.id as id0_0_, person0_.name as name0_0_
from Person person0_ where person0_.id=?

**org.hibernate.NonUniqueObjectException: a different object with the same identifier value was
already associated with the session: [entity.Person#1]**

Exception




Merge – Avoids the Exception





Merge - Misunderstood

- Merge does not behave like `save()`, `update()`, and `saveOrUpdate()`
 - The object you pass in to merge never becomes part of the persistence context (is not managed)
 - Instead merge returns an different object representing the same entity (which *is* managed)
- If you continue working with the original object you can run into unexpected problems, implicit updates to it are not persisted





Correct use of Merge



Correct Use of merge() – Always use the return value

```
Person p = new Person("Frank Brown");

try {
    session = sessionFactory.openSession();
    tx = session.beginTransaction();

    p = (Person)session.merge(p);
    p.setName("Upd. Frank Brown");

    tx.commit();
} catch (HibernateExc
tx.rollback();
e.printStackTrace();
} finally {
    if (session != null) {
        session.close();
    }
}
```

P is set to
the return
of merge()

Update will be
committed



Return is
lost, P is
still the
detached object

Incorrect Use of merge()



```
Person p = new Person("Frank Brown");

try {
    session = sessionFactory.openSession();
    tx = session.beginTransaction();

    session.merge(p);
    p.setName("Upd. Frank Brown");

    tx.commit();
} catch (HibernateException e) {
    tx.rollback();
    e.printStackTrace();
} finally {
    if (session != null) {
        session.close();
    }
}
```

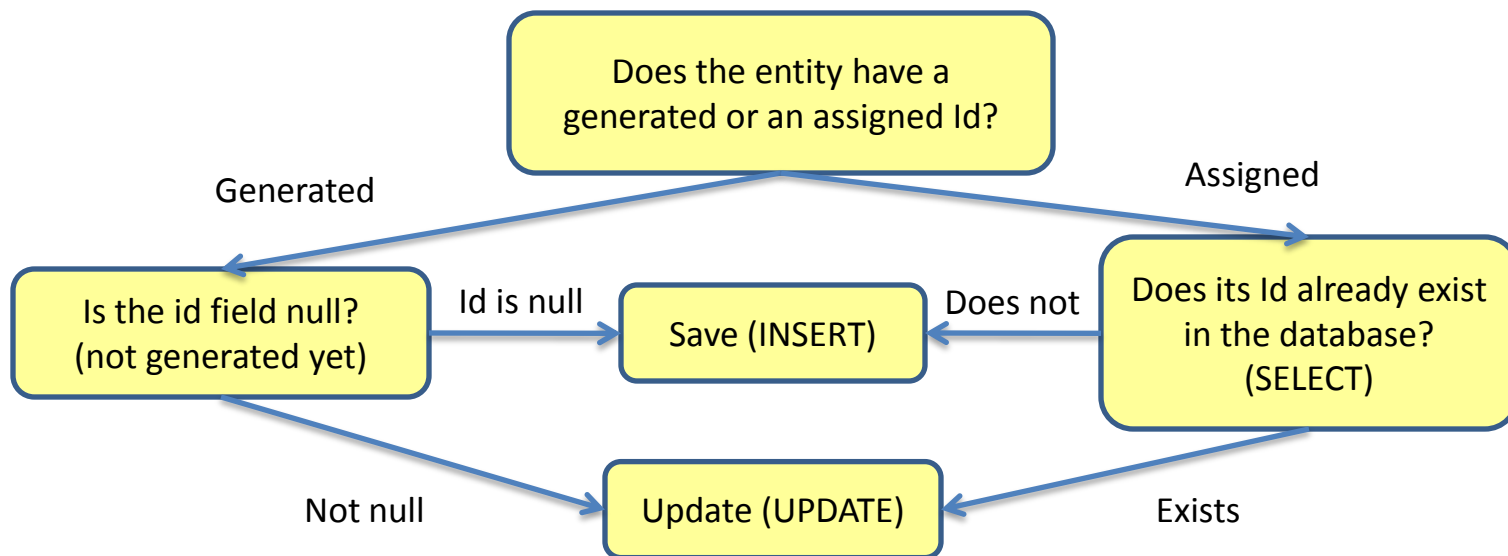
Update will
never get
committed





saveOrUpdate()

- Hibernate checks if it needs to INSERT or UPDATE
 - If needed using a SELECT to find if it's already in the db





Persistence API:

REMOVING ENTITY OBJECTS



About: delete()

- Can be used to **remove an entity from the db**
 - Only managed entities can be removed
- Will only work inside a transaction
- On `transaction.commit()` a SQL DELETE is used to delete the entity from the database

```
try {  
    session = sessionFactory.openSession();  
    tx = session.beginTransaction();  
  
    Person p1 = (Person) session.load(Person.class, 1L);  
    session.delete(p1);  
  
    tx.commit();  
} catch (HibernateException e) {  
    tx.rollback();  
    e.printStackTrace();  
}
```

Delete entity

Load entity so that is managed in the persistence context

SQL DELETE is not executed until the transaction commits



Persistence API:

PERSISTENCE CACHE



When does Hibernate access the DB?

```
session = sessionFactory.openSession();  
tx = session.beginTransaction();
```

```
Employee fmemmployee = new Employee("Frank", "Miller");  
System.out.println("1");  
session.persist(fmemmployee);  
System.out.println("2");  
  
tx.commit();
```

Hibernate will access the database during the persist() to retrieve the unique key

```
@Entity  
public class Employee {  
  
    @Id  
    @GeneratedValue  
    private int id;  
    private String firstname;  
    private String lastname;  
    ...  
}
```

```
1  
Hibernate: /* insert Employee */ insert into Employee (id, lastname) values (null, ?)  
Hibernate: call identity()  
2
```



When does Hibernate access the DB?



```
session = sessionFactory.openSession();  
tx = session.beginTransaction();
```

```
Employee fmemmployee = new Employee("Frank", "Miller");  
System.out.println("1");  
session.persist(fmemmployee);  
System.out.println("2");
```

```
tx.commit();  
System.out.println("3");
```

Hibernate will access
the database during the
commit()

```
@Entity
```

```
public class Employee {  
  
    @Id  
    private int id;  
    private String firstname;  
    private String lastname;  
    ...  
}
```

Application
generates the id

```
1  
2  
Hibernate: /* insert Employee */ insert into Employee (firstname, lastname, id) values  
  (?, ?, ?)  
3
```



When does Hibernate access the DB?

```
session = sessionFactory.openSession();  
tx = session.beginTransaction();
```

```
Employee fmemmployee = new Employee("Frank", "Miller");  
System.out.println("1");  
session.persist(fmemmployee);  
System.out.println("2");  
fmemmployee.setFirstname("John");  
System.out.println("3");  
fmemmployee.setLastname("Doe");  
System.out.println("4");
```

Hibernate will access the database during the persist() to retrieve the unique key

```
tx.commit();  
System.out.println("5");
```

Hibernate will access the database during the commit() to update the state of the changed employee

```
@Entity  
public class Employee {  
  
    @Id  
    @GeneratedValue  
    private int id;  
    private String firstname;  
    private String lastname;  
    ...  
}
```

```
1  
Hibernate: /* insert Employee */ insert into Employee (id, firstname, lastname) values  
  (null, ?, ?)  
Hibernate: call identity()  
2  
3  
4  
Hibernate: /* update Employee */ update Employee set firstname=?, lastname=? where id=?  
5
```



When does Hibernate access the DB?

```
session = sessionFactory.openSession();
tx = session.beginTransaction();

Employee fmemmployee = new Employee("Frank", "Miller");
System.out.println("1");
session.persist(fmemmployee);
System.out.println("2");
fmemmployee.setFirstname("John");
System.out.println("3");
fmemmployee.setLastname("Doe");
System.out.println("4");
Employee jdemmployee= (Employee)session.get(Employee.class, fmemmployee.getId());
System.out.println("5");

tx.commit();
System.out.println("6");
```

No database access,
object is already in the
persistence context

```
1
Hibernate: /* insert Employee */ insert into Employee (id, firstname, lastname) values
(null, ?, ?)
Hibernate: call identity()
2
3
4
5
Hibernate: /* update Employee */ update Employee set firstname=?, lastname=? where id=?
6
```

When does Hibernate access the DB?

```
session = sessionFactory.openSession();  
tx = session.beginTransaction();
```

```
Employee fmemmployee = new Employee("Frank", "Miller");  
System.out.println("1");  
session.persist(fmemmployee);  
System.out.println("2");  
fmemmployee.setFirstname("John");  
System.out.println("3");  
fmemmployee.setLastname("Doe");  
System.out.println("4");  
Query query= session.createQuery("from Employee");  
Collection<Employee> employeeelist= query.list();  
System.out.println("5");  
  
tx.commit();  
System.out.println("6");
```

```
@Entity  
public class Employee {  
  
    @Id  
    @GeneratedValue  
    private int id;  
    private String firstname;  
    private String lastname;  
    ...  
}
```

Query on Employee, and we have a dirty Employee in the persistence context. First update the Employee in the database, and then execute the query

```
1  
Hibernate: /* insert Employee */ insert into Employee (id, firstname, lastname) values  
  (null, ?, ?)  
Hibernate: call identity()  
2  
3  
4  
Hibernate: /* update Employee */ update Employee set firstname=?, lastname=? where id=?  
Hibernate: /* from Employee */ select employee0_.id as id0_, employee0_.firstname as  
  firstname0_, employee0_.lastname as lastname0_ from Employee employee0_  
5  
6
```



flush()

- Flushes any changes held in the session persistence context to the database
- Any changes made by an implicit or explicit update(), objects marked for deletion with delete() or for persistence with persist()
- All these changes can be committed to the database immediately **without having to wait for the transaction to commit**

```
session.flush();
```



flush()

```
Session session = sessionFactory.openSession();
Transaction tx = session.beginTransaction();
```

```
APerson p1 = new APerson("Frank Brown", 1);
APerson p2 = new APerson("John Doe", 2);
APerson p3 = new APerson("Mary Smith", 3);
session.persist(p1);
session.persist(p2);
session.persist(p3);
```

Persist() does not insert entities with assigned ids right away

```
p1.setName("Frank Updated");
p2.setName("John Updated");
session.delete(p3);
```

Creates changes to the objects in the persistence cache

```
System.out.println("About to flush()");
session.flush();
System.out.println("Changes flushed");
```

flush() everything

```
tx.commit();
session.close();
```

No changes left to commit()

Flush() commits all changes held in cache

About to flush()

```
Hibernate: /* insert entity.APerson */ insert into APerson (name, id) values (?, ?)
Hibernate: /* insert entity.APerson */ insert into APerson (name, id) values (?, ?)
Hibernate: /* insert entity.APerson */ insert into APerson (name, id) values (?, ?)
Hibernate: /* update entity.APerson */ update APerson set name=? where id=?
Hibernate: /* update entity.APerson */ update APerson set name=? where id=?
Hibernate: /* delete entity.APerson */ delete from APerson where id=?
Changes flushed
```

Person with assigned identifier

```
@Entity
public class APerson {
    @Id
    private long id;
    private String name;
    ...
}
```

DB Result:


```
mysql> select * from person;
+----+-----+
| id | name          |
+----+-----+
|  1 | Frank Updated |
|  2 | John Updated  |
+----+-----+
```

Hibernate Output:



When are changes Flushed?

- There are actually 3 different scenarios under which Hibernate will commit all changes
 1. When an explicit `session.flush()` call is made
 2. When a transaction is committed
 3. Inside a transaction, before a query to the database is executed





refresh()

- The entity object will immediately be refreshed with the **state currently held in the database**
- Can be used to undo updates

```
session = sessionFactory.openSession();  
tx = session.beginTransaction();  
  
session.refresh(p1);  
  
tx.commit();  
session.close();
```

Although a transaction is not required, it is recommended to always use a transaction

Immediately refreshes with state from the database, re-attaches if detached

Stale Cache – a Reason to Refresh

```
try {
    session = sessionFactory.openSession();
    tx = session.beginTransaction();

    Person p1 = (Person) session.get(Person.class, 1L);
    System.out.println(p1.getName());

    Thread.sleep(1000 * 60 * 5);
    System.out.println("Slept for 5 minutes");

    p1 = (Person) session.get(Person.class, 1L);
    System.out.println(p1.getName());

    session.refresh(p1);
    System.out.println(p1.getName());

    tx.commit();
} catch (Exception e) {
    ...
}
```

Retrieve a person entity and print its name (also puts it in the persistence cache)

Try to retrieve the same entity again 5 minutes later

Because of cache get() does not check the database!

Specifically refresh() to check the database for updates

Hibernate Output:

```
Hibernate: /* load entity.Person */ select person0_.id as id0_0_,
person0_.name as name0_0_ from Person person0_ where person0_.id=?
Frank Brown
Slept for 5 minutes
Frank Brown
Hibernate: /* load entity.Person */ select person0_.id as id0_0_,
person0_.name as name0_0_ from Person person0_ where person0_.id=?
Frank Upd.
```

refresh() checks the database and finds that it has been updated!

No db check, just cache



Entity Manager API

Method	Description
<code>clear()</code>	Empties persistence context, managed entities become detached
<code>close()</code>	Closes the entity manager, managed entities become detached
<code>contains()</code>	Checks if the instance is managed by this entity manager
<code>detach()</code>	Detaches the given entity
<code>find()</code>	Retrieves and entity from the db by primary key
<code>flush()</code>	Flushes all changes to the db
<code>getReference()</code>	Get and instance by Id, state may be lazily fetched
<code>merge()</code>	Merge the state of the given entity into the persistence context
<code>persist()</code>	Make the instance managed and persistent
<code>refresh()</code>	Update the state of the given entity from the db
<code>remove()</code>	Delete the given entity from the db

<http://docs.oracle.com/javaee/6/api/javax/persistence/EntityManager.html>



Active Learning

- What are the four different states of the entity lifecycle?
- What are the similarities and the differences between `saveOrUpdate()` and `merge()`?



Module Summary

- In this module we discussed the object persistence lifecycle and how to manage it
- We presented the different methods Hibernate provides on the Session object
- Although several methods provide similar functionality there are subtle differences
- These differences are important for enterprise applications and should be noted with care
- We also demonstrated Hibernates caching behavior and how it can be manipulated