



Module 09: Hibernate Optimization

CS544: Enterprise Architecture



Data Access Optimization

- By changing when and How Hibernate retrieves data you can optimize its performance
- By default Hibernate tries to not load data that you might not need – mostly uses lazy loading
 - This gives good baseline performance
 - Is not always the most optimal strategy





When and How

- When data is retrieved is closely related to how data is retrieved
 - Hibernate provides several fetching strategies

Strategy	Description
Lazy Loading	Loads data only when needed, using a select for each item
Batch Fetching	Loads more of the same entity / collection when one is needed
Sub Select	Loads all collections for retrieved entities when one is needed
Eager Join	Immediately loads related data, using joins if possible else selects
Join Fetch Query	A query that also loads related entities or collections





Hibernate Optimization:

PERFORMANCE PROBLEMS



Slow? → What to Look For



- There are two main problem categories:
 - Many selects to get similar, or closely related data
 - These selects can probably be combined
 - The N + 1 problem is an example of this
 - Caused by inappropriate lazy loading of data
 - Complex queries that use many joins
 - May be more efficient to use several simple queries
 - Cartesian Product problem is an example of this
 - Caused by incorrect (over) optimization



The N + 1 Problem

4 Customers each with
their own salesrep

```
session = sessionFactory.openSession();
tx = session.beginTransaction();

Customer cust1 = new Customer("Frank", "Brown");
Customer cust2 = new Customer("Jane", "Terrien");
Customer cust3 = new Customer("John", "Doe");
Customer cust4 = new Customer("Carol", "Reno");
cust1.setSalesRep(new SalesRep("John Willis"));
cust2.setSalesRep(new SalesRep("Mary Long"));
cust3.setSalesRep(new SalesRep("Ted Walker"));
cust4.setSalesRep(new SalesRep("Keith Rogers"));

session.persist(cust1);
session.persist(cust2);
session.persist(cust3);
session.persist(cust4);

tx.commit();
```

```
List<Customer> customers =
    session.createQuery("from Customer").list();
SalesRep salesrep = null;
for (Customer customer : customers) {
    salesrep = customer.getSalesRep();
    // do something with the salesrep
    salesrep.getName();
}
```


Retrieve customers and then
work with related salesrep

Gets the customers (1 query), then
executes another query for each
salesrep (N queries). Total N + 1

```
Hibernate:
select
    customer0_.id as id0_,
    customer0_.firstname as firstname0_,
    customer0_.lastname as lastname0_,
    customer0_.salesRep_id as salesRep4_0_
from
    Customer customer0_
Hibernate:
select
    salesrep0_.id as id1_0_,
    salesrep0_.name as name1_0_
from
    SalesRep salesrep0_
where
    salesrep0_.id=?
Hibernate:
select
    salesrep0_.id as id1_0_,
    salesrep0_.name as name1_0_
from
    SalesRep salesrep0_
where
    salesrep0_.id=?
Hibernate:
select
    salesrep0_.id as id1_0_,
    salesrep0_.name as name1_0_
from
    SalesRep salesrep0_
where
    salesrep0_.id=?
```



N + 1 with References (to-one)





N + 1 with Collections

2 Sales Reps, each with a collection of customers

```
session = sessionFactory.openSession();
tx = session.beginTransaction();

SalesRep sr1 = new SalesRep("John Willis");
SalesRep sr2 = new SalesRep("Mary Long");

sr1.addCustomer(new Customer("Frank", "Brown"));
sr1.addCustomer(new Customer("Jane", "Terrien"));
sr2.addCustomer(new Customer("John", "Doe"));
sr2.addCustomer(new Customer("Carol", "Reno"));

session.persist(sr1);
session.persist(sr2);

tx.commit();
```

```
List<SalesRep> salesReps =
    session.createQuery("from SalesRep").list();
for (SalesRep s : salesReps) {
    Set<Customer> customers = s.getCustomers();
    for (Customer c : customers) {
        // do something with the customer
    }
}
```

Retrieve sales reps, and then work with related customers

Gets the sales reps (1 query), and then executes another query for each sales rep (N queries). Total N + 1 queries

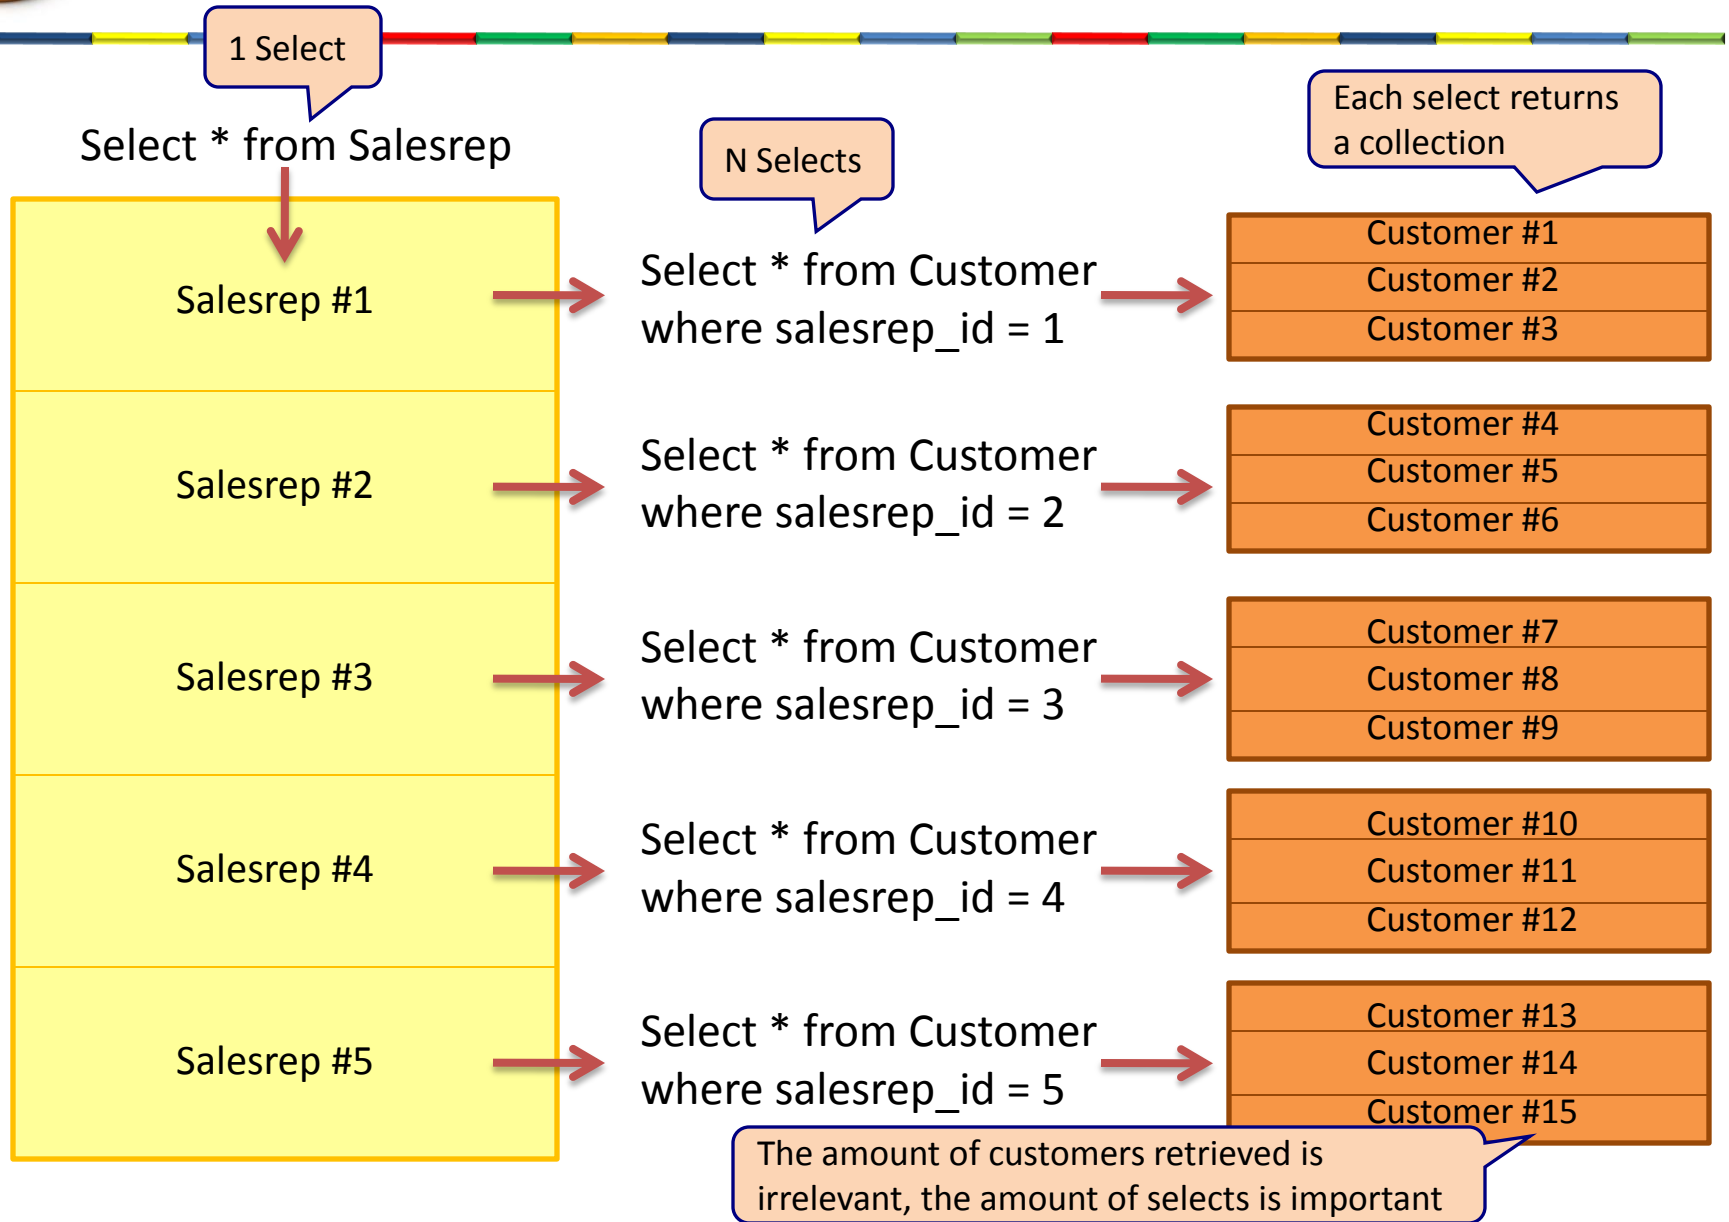
```
Hibernate:
select
  salesrep0_.id as id1_,
  salesrep0_.name as name1_
from
  SalesRep salesrep0_

Hibernate:
select
  customers0_.salesRep_id as salesRep4_1_,
  customers0_.id as id1_,
  customers0_.id as id0_0_,
  customers0_.firstname as firstname0_0_,
  customers0_.lastname as lastname0_0_,
  customers0_.salesRep_id as salesRep4_0_0_
from
  Customer customers0_
where
  customers0_.salesRep_id=?

Hibernate:
select
  customers0_.salesRep_id as salesRep4_1_,
  customers0_.id as id1_,
  customers0_.id as id0_0_,
  customers0_.firstname as firstname0_0_,
  customers0_.lastname as lastname0_0_,
  customers0_.salesRep_id as salesRep4_0_0_
from
  Customer customers0_
where
  customers0_.salesRep_id=?
```




N + 1 with Collections (to-many)





Cartesian Product Problem

```
Customer cust1 = new Customer("Frank", "Brown");
Customer cust2 = new Customer("Jane", "Terrien");
Customer cust3 = new Customer("John", "Doe");

cust1.addBook(new Book("Harry Potter and the Deathly Hallows"));
cust1.addBook(new Book("Unseen Academicals (Discworld)"));
cust1.addBook(new Book("The Color of Magic (Discworld)"));
cust1.addMovie(new Movie("Shrek"));
cust1.addMovie(new Movie("WALL-E"));
cust1.addMovie(new Movie("Howls Moving Castle"));

cust2.addBook(new Book("Twilight (The Twilight Saga, Book1)"));

cust3.addMovie(new Movie("Forgetting Sarah Marshall"));
```

Customers have a set of books, and a set of movies that they like.

First customer has 3 books and 3 movies, second customer has a single book, third customer has a single movie

Retrieve customers, and also try to (eager) fetch the book and movie collections for the customers

Hibernate:
select

```
customer0_.firstname as firstname0_0_,
customer0_.lastname as lastname0_0_,
books1_.title as title1_1_,
movies2_.title as title2_2_
from
  Customer customer0_
left outer join
  Book books1_
    on customer0_.id=books1_.customer_id
left outer join
  Movie movies2_
    on customer0_.id=movies2_.customer_id
```

FIRSTNAME0_0_	LASTNAME0_0_	TITLE1_1_	TITLE2_2_
Frank	Brown	Unseen Academicals (Discworld)	WALL-E
Frank	Brown	Unseen Academicals (Discworld)	Shrek
Frank	Brown	Unseen Academicals (Discworld)	Howls Moving Castle
Frank	Brown	The Color of Magic (Discworld)	WALL-E
Frank	Brown	The Color of Magic (Discworld)	Shrek
Frank	Brown	The Color of Magic (Discworld)	Howls Moving Castle
Frank	Brown	Harry Potter and the Deathly Hallows	WALL-E
Frank	Brown	Harry Potter and the Deathly Hallows	Shrek
Frank	Brown	Harry Potter and the Deathly Hallows	Howls Moving Castle
Jane	Terrien	Twilight (The Twilight Saga, Book1)	[null]
John	Doe	[null]	Forgetting Sarah Marshall

Outer Joining two or more collections creates many redundant rows

Row count per customers = related books * related movies



Cartesian Product

- Joining two collections creates: $R \times N \times M$
 - Creating a very in-efficient resultset

Frank Brown ✓	Discworld ✓	Pixar ✓
Frank Brown	Discworld	Dream Works ✓
Frank Brown	Discworld	Studio Ghibli ✓
Frank Brown	Harry Potter ✓	Pixar
Frank Brown	Harry Potter	Dream Works
Frank Brown	Harry Potter	Studio Ghibli
Frank Brown	Twilight ✓	Pixar
Frank Brown	Twilight	Dream Works
Frank Brown	Twilight	Studio Ghibli

9 rows, 3 columns to give 7 pieces of data



Hibernate Optimization:

STRATEGY: LAZY LOADING



Lazy Loading

- Lazy loading can be specified for:
 - **Object References**
 - one-to-one and many-to-one associations
 - Generally don't default to lazy loading (good thing)
 - **Collections**
 - one-to-many and many-to-many associations
 - Have the option to use 'extra-lazy' loading
 - **Large Properties**
 - CLOBs and BLOBs, e.g. large texts or image data
 - Need byte code instrumentation to use lazy loading



Annotations

- The JPA specifies that both `@ManyToOne` and `@OneToOne` default to eager loading

```
@Entity
public class Customer {
    @Id
    @GeneratedValue
    private int id;
    private String firstname;
    private String lastname;

    @OneToOne(cascade=CascadeType.PERSIST)
    private Address address;

    @ManyToOne
    private SalesRep salesRep;
```

```
Customer cust1 = (Customer)
    session.get(Customer.class, 1);
```

Hibernate retrieves the customer,
the address, and the salesrep

```
Hibernate:
select
    customer0_.id as id0_2_,
    customer0_.address_id as address4_0_2_,
    customer0_.firstname as firstname0_2_,
    customer0_.lastname as lastname0_2_,
    customer0_.salesRep_id as salesRep5_0_2_,
    address1_.id as id1_0_,
    address1_.apt as apt1_0_,
    address1_.city as city1_0_,
    address1_.state as state1_0_,
    address1_.street as street1_0_,
    address1_.zip as zip1_0_,
    salesrep2_.id as id3_1_,
    salesrep2_.name as name3_1_
from
    Customer customer0_
left outer join
    Address address1_
        on customer0_.address_id=address1_.id
left outer join
    SalesRep salesrep2_
        on customer0_.salesRep_id=salesrep2_.id
where
    customer0_.address_id=?
```



Specifying Lazy

- Although it is possible to set @OneToOne and @ManyToOne to FetchType.LAZY

```
@Entity
public class Customer {
    @Id
    @GeneratedValue
    private int id;
    private String firstname;
    private String lastname;
```

FetchType.LAZY

```
@OneToOne(fetch = FetchType.LAZY,
           cascade=CascadeType.PERSIST)
private Address address;
```

FetchType.LAZY

```
@ManyToOne(fetch = FetchType.LAZY)
private SalesRep salesRep;
...
```

```
Customer cust1 = (Customer)
    session.get(Customer.class, 1);
```

Hibernate again only retrieves the customer object

Hibernate:
select

```
    customer0_.id as id0_0_,
    customer0_.address_id as address4_0_0_,
    customer0_.firstname as firstname0_0_,
    customer0_.lastname as lastname0_0_,
    customer0_.salesRep_id as salesRep5_0_0_
from
    Customer customer0_
where
    customer0_.id=?
```



XML

Can only be changed with
byte-code instrumentation

- `<one-to-one>` defaults to eager loading ⚡
- `<many-to-one>` defaults to lazy loading

```
public class Customer {  
    private int id;  
    private String firstname;  
    private String lastname;  
    private Address address;  
    private SalesRep salesRep;  
    ...  
}
```

```
Customer cust1 = (Customer)  
    session.get(Customer.class, 1);
```

Eagerly loads `<one-to-one>`
using a left outer join

Hibernate:
select

```
    customer0_.id as id0_1_,  
    customer0_.firstname as firstname0_1_,  
    customer0_.lastname as lastname0_1_,  
    customer0_.salesRep as salesRep0_1_,  
    address1_.id as id1_0_,  
    address1_.street as street1_0_,  
    address1_.apt as apt1_0_,  
    address1_.city as city1_0_,  
    address1_.state as state1_0_,  
    address1_.zip as zip1_0_,  
    address1_.customer as customer1_0_  
from  
    Customer customer0_  
left outer join  
    Address address1_  
        on customer0_.id=address1_.id  
where  
    customer0_.id=?
```

```
<hibernate-mapping package="when.objRefs">  
    <class name="Customer">  
        <id name="id">  
            <generator class="native" />  
        </id>  
        <property name="firstname" />  
        <property name="lastname" />  
        <one-to-one name="address" cascade="persist" />  
        <many-to-one name="salesRep" />  
    </class>  
</hibernate-mapping>
```




Collections

- By default the entire collection is retrieved when `.size()`, `.isEmpty()`, or `.contains()` is used
 - Good for small collections, bad for large collections

```
@Entity
public class Customer {
    @Id
    @GeneratedValue
    private int id;
    private String firstname;
    private String lastname;

    @OneToMany(mappedBy = "customer",
        cascade = CascadeType.PERSIST)
    private Set<CreditCard> creditCards
        = new HashSet<CreditCard>();

    ...
}
```

Customer with
a collection of
Credit Cards

Check credit card
collection size

```
customer.getCreditCards().size();
```

Retrieves all
credit cards

```
Hibernate:
select
    creditcard0_.customer_id as customer5_1_,
    creditcard0_.id as id1_,
    creditcard0_.id as id1_0_,
    creditcard0_.customer_id as
        customer5_1_0_,
    creditcard0_.expiration as expiration1_0_,
    creditcard0_.name as name1_0_,
    creditcard0_.number as number1_0_
from
    CreditCard creditcard0_
where
    creditcard0_.customer_id=?
```





Extra Lazy Collections

- Setting the collection to **Extra Lazy** solves this problem for large collections

Extra Lazy Collection
using annotations

```
@Entity
public class Customer {
    @Id
    @GeneratedValue
    private int id;
    private String firstname;
    private String lastname;

    @OneToMany(mappedBy = "customer", cascade = CascadeType.PERSIST)
    @org.hibernate.annotations.LazyCollection(
        org.hibernate.annotations.LazyCollectionOption.EXTRA
    )
    private Set<CreditCard> creditCards = new HashSet<CreditCard>();

    ...
}
```

```
customer.getCreditCards().size();
```

```
<hibernate-mapping package="entities">
  <class name="Customer">
    <id name="id">
      <generator class="native" />
    </id>
    <property name="firstname" />
    <property name="lastname" />
    <set name="creditCards" lazy="extra"
      inverse="true" cascade="persist">
      <key column="customer" />
      <one-to-many class="CreditCard" />
    </set>
  </class>
</hibernate-mapping>
```

Extra Lazy Collection
using XML mapping

Only retrieves size

```
Hibernate:
select
  count(id)
from
  CreditCard
where
  customer_id =?
```



Large Properties

- Certain Properties may be so large that you only want to load them when really necessary
 - Lazy loading of properties is only available with byte-code instrumentation

```
public class Book {  
    private String isbn;  
    private String title;  
    private String author;  
    private java.sql.Clob summary;  
    private java.sql.Blob cover;  
    ...  
}
```

```
<hibernate-mapping package="when.lazyprops">  
    <class name="Book">  
        <id name="isbn" />  
        <property name="title" />  
        <property name="author" />  
        <property name="summary" type="clob" lazy="true" />  
        <property name="cover" type="blob" lazy="true" />  
    </class>  
</hibernate-mapping>
```

Without byte-code
instr. lazy=true
doesn't do anything

```
Book b = (Book) session.get(Book.class, "978-0545139700");
```

Summary and cover
are not loaded (lazy)

Hibernate:

```
select  
    book0_.isbn as isbn0_0_,  
    book0_.title as title0_0_,  
    book0_.author as author0_0_  
from  
    Book book0_  
where  
    book0_.isbn=?
```



Byte-Code Instrumentation Ant File

```
<?xml version="1.0" encoding="UTF-8"?>
<project name="ByteCodeInstrument" default="instrument">
  <description>Byte Code instrument example</description>
  <property name="src" location="src" />
  <property name="build" location="bin" />

  <target name="compile">
    <javac srcdir="${src}" destdir="${build}" />
  </target>

  <target name="instrument" depends="compile">
    <taskdef name="instrument"
      classname="org.hibernate.tool.instrument.cglib.InstrumentTask">
      <classpath>
        <fileset dir="c:/hibernatetraining/libraries/">
          <include name="**/*.jar" />
        </fileset>
      </classpath>
    </taskdef>
    <instrument verbose="true">
      <fileset dir="${build}/when/properties/">
        <include name="**/*.class" />
      </fileset>
    </instrument>
  </target>
</project>
```

Code needs to be
compiled before it
can be instrumented

Requires the hibernate libraries

Location of the files that
need to be instrumented



Annotations – Lazy Properties

- Requires **property access** for lazy loading

```
Book b = (Book)session.get(Book.class, "978-0545139700");  
System.out.println(b.getTitle());
```

```
java.sql.Clob sumData = b.getSummary();  
int length = (int)sumData.length();  
System.out.println(sumData.getSubString(1, length));
```

```
@Entity  
public class Book {  
    ...  
  
    @Id  
    public String getIsbn() { return isbn; }  
    public String getTitle() { return title; }  
    public String getAuthor() { return author; }  
  
    @Basic(fetch=FetchType.LAZY)  
    public java.sql.Clob getSummary() {  
        return summary;  
    }  
    @Basic(fetch=FetchType.LAZY)  
    public java.sql.Blob getCover() {  
        return cover;  
    }  
    ...  
}
```

Annotations on getter methods instead of fields for property access

Both Summary and Cover are loaded

Hibernate:

```
select  
    book0_.isbn as isbn0_0_,  
    book0_.title as title0_0_,  
    book0_.author as author0_0_  
from  
    Book book0_  
where  
    book0_.isbn=?  
Harry Potter and the Deathly Hallows
```

Hibernate:

```
select  
    book_.summary as summary0_,  
    book_.cover as cover0_  
from  
    Book book_  
where  
    book_.isbn=?  
Readers beware. The brilliant,  
breathtaking conclusion to J.K.  
Rowling's spellbinding series is not for  
the faint of heart
```

Only loads summary when needed



Lazy Loading

- Lazy loading gives good baseline performance
 - But it is also the cause of **N + 1 problems**
 - Can be slow because of too many small selects
- **Solution:** fetch data more eagerly
 - To which strategy you switch should depend on the circumstance surrounding the association





Hibernate Optimization:

STRATEGY: EAGER JOIN



Eager Join

- The eager join strategy is a combination of (non lazy) **Eager** behavior and **Join Fetching**
 - Using XML these are two separate things
 - Using annotations they can not be separated
- Similar but different:
 - Eager specifies that a relation has to be loaded right away if it hasn't been loaded already
 - Join Fetching retrieves a related entity or collection at the same time using a SQL Join
 - **Joins are always eager, eager is not always with joins**
 - Will complain when multiple collections are eager



Eager - XML

Hibernate:

select

```
customer0_.id as id0_0_,
customer0_.firstname as firstname0_0_,
customer0_.lastname as lastname0_0_,
customer0_.salesRep as salesRep0_0_
```

from

```
Customer customer0_
```

where

```
customer0_.id=?
```

Hibernate:

select

```
salesrep0_.id as id1_0_,
salesrep0_.name as name1_0_
```

from

```
SalesRep salesrep0_
```

where

```
salesrep0_.id=?
```

Hibernate:

select

```
creditcard0_.customer as customer1_,
creditcard0_.id as id1_,
creditcard0_.id as id2_0_,
creditcard0_.number as number2_0_,
creditcard0_.name as name2_0_,
creditcard0_.expiration as expiration2_0_,
creditcard0_.customer as customer2_0_
```

from


```
CreditCard creditcard0_
```

where

```
creditcard0_.customer=?
```

Eager Many-to-one and
One-To-Many relation

```
<hibernate-mapping package="eager.xml">
  <class name="Customer">
    <id name="id">
      <generator class="native">
        </generator>
      </id>
    <property name="firstname" />
    <property name="lastname" />
    <many-to-one name="salesRep" lazy="false" />
    <set name="creditCards" inverse="true"
      cascade="persist" lazy="false">
      <key column="customer" />
      <one-to-many class="CreditCard" />
    </set>
  </class>
</hibernate-mapping>
```



```
session.get(Customer.class, 1);
```

Get() a single customer creates 3 selects



Join Fetching – XML

```
<hibernate-mapping package="how.always.join">
  <class name="Customer">
    <id name="id">
      <generator class="native" />
    </id>
    <property name="firstname" />
    <property name="lastname" />
    <many-to-one name="salesRep" fetch="join" />
    <set name="creditCards" inverse="true"
      cascade="persist" fetch="join">
      <key column="customer" />
      <one-to-many class="CreditCard" />
    </set>
  </class>
</hibernate-mapping>
```

Join Fetching for Many-To-One
and One-To-Many

Joined collections can lead
to a Cartesian product

```
session.get(Customer.class, 1);
```

Get() loads three entities using a single select

date:
select

```
customer0_.id as id0_2_,
customer0_.firstname as firstname0_2_,
customer0_.lastname as lastname0_2_,
customer0_.salesRep_id as salesRep4_0_2_,
creditcard1_.customer_id as customer5_4_,
creditcard1_.id as id4_,
creditcard1_.id as id2_0_,
creditcard1_.customer_id as customer5_2_0_,
creditcard1_.expiration as expiration2_0_,
creditcard1_.name as name2_0_,
creditcard1_.number as number2_0_,
salesrep2_.id as id1_1_,
salesrep2_.name as name1_1_
```

from

```
Customer customer0_
```

left outer join

```
CreditCard creditcard1_
```

```
on customer0_.id=creditcard1_.customer_id
```

left outer join

```
SalesRep salesrep2_
```

```
on customer0_.salesRep_id=salesrep2_.id
```

where

```
customer0_.id=?
```



Eager and Join - Annotations

@Entity

```
public class Customer {  
    @Id  
    @GeneratedValue  
    private int id;  
    private String firstname;  
    private String lastname;  
  
    @ManyToOne(fetch=FetchType.EAGER)  
    private SalesRep salesRep;  
  
    @OneToMany(fetch=FetchType.EAGER,  
        mappedBy="customer",  
        cascade=CascadeType.PERSIST)  
    private Set<CreditCard> creditCards =  
        new HashSet<CreditCard>();  
    ...  
}
```

Specify eager and join
using FetchType.EAGER

Hibernate:

```
select  
    customer0_.id as id0_2_,  
    customer0_.firstname as firstname0_2_,  
    customer0_.lastname as lastname0_2_,  
    customer0_.salesRep_id as salesRep4_0_2_,  
    creditcard1_.customer_id as customer5_4_,  
    creditcard1_.id as id4_,  
    creditcard1_.id as id2_0_,  
    creditcard1_.customer_id as customer5_2_0_,  
    creditcard1_.expiration as expiration2_0_,  
    creditcard1_.name as name2_0_,  
    creditcard1_.number as number2_0_,  
    salesrep2_.id as id1_1_,  
    salesrep2_.name as name1_1_  
from  
    Customer customer0_  
left outer join  
    CreditCard creditcard1_  
        on customer0_.id=creditcard1_.customer_id  
left outer join  
    SalesRep salesrep2_  
        on customer0_.salesRep_id=salesrep2_.id  
where  
    customer0_.id=?
```

Joined collections are
dangerous

```
session.get(Customer.class, 1);
```

All three entities are loaded
using a single joined select



No Problem?

- So far combining Eager and Join in the same specification doesn't seem like a problem
 - Eager without join seems silly, just following the references would give the same result
 - Joins are always eager anyway
 - Surely there is no problem in combining these!



Eager and Queries

- Queries only fetch what is specified in the query
 - Mapping based optimizations (what / how to retrieve) are ignored when executing a query
- Eager mapping specifies **when** not what or how
 - After the query has executed, hibernate will fulfill the eager specification **using extra selects**
 - In other words, a **guaranteed N+1** for all eager associations not yet retrieved by the query
 - Both for references and collections



Annotations Eager Join & Query

@Entity

```
public class Customer {  
    @Id  
    @GeneratedValue  
    private int id;  
    private String firstname;  
    private String lastname;  
  
    @ManyToOne(fetch = FetchType.EAGER)  
    private SalesRep salesRep;  
  
    @OneToMany(fetch=FetchType.EAGER,  
        mappedBy="customer",  
        cascade=CascadeType.PERSIST)  
    private Set<CreditCard> creditCards =  
        new HashSet<CreditCard>();  
  
    ...  
}
```

Creates an N+1 problems
without even having a loop!

```
Query query =  
    session.createQuery("from Customer");  
List<Customer> customers = query.list();
```

Hibernate:

```
select  
    customer0_.id as id0_,  
    customer0_.firstname as firstname0_,  
    customer0_.lastname as lastname0_,  
    customer0_.salesRep_id as salesRep4_0_  
from  
    Customer customer0_
```

N selects, one for each
customer retrieved

Hibernate:

```
select  
    salesrep0_.id as id1_0_,  
    salesrep0_.name as name1_0_  
from  
    SalesRep salesrep0_  
where  
    salesrep0_.id=?
```

N selects, one for each
customer retrieved

Hibernate:

```
select  
    creditcard0_.customer_id as customer5_1_,  
    creditcard0_.id as id1_,  
    creditcard0_.id as id2_0_,  
    creditcard0_.customer_id as customer5_2_0_,  
    creditcard0_.expiration as expiration2_0_,  
    creditcard0_.name as name2_0_,  
    creditcard0_.number as number2_0_  
from  
    CreditCard creditcard0_  
where  
    creditcard0_.customer_id=?
```



XML Join & Query

```
<hibernate-mapping package="how.always.join">
  <class name="Customer">
    <id name="id">
      <generator class="native" />
    </id>
    <property name="firstname" />
    <property name="lastname" />
    <many-to-one name="salesRep" fetch="join" />
    <set name="creditCards" cascade="persist"
      inverse="true" fetch="join" >
      <key column="customer" />
      <one-to-many class="CreditCard" />
    </set>
  </class>
</hibernate-mapping>
```

```
Query query =
    session.createQuery("from Customer");
List<Customer> customers = query.list();
```

XML join mappings default to lazy="true" thereby not causing the N+1 problems

Unfortunately no such option exists for annotations

Hibernate:
select

```
customer0_.id as id0_,
customer0_.firstname as firstname0_,
customer0_.lastname as lastname0_,
customer0_.salesRep as salesRep0_
from
    Customer customer0_
```





Beware: Eager Join

- Because of the problems eager join can create we mostly recommend against using it
 - There are still valid reasons why you might want to use join fetching, especially with XML
 - Just beware that eager can cause **N+1 problems**





Hibernate Optimization:

STRATEGY: JOIN FETCH QUERY



Join Fetch Query

- A Join Fetch Query is the most flexible strategy
 - Other strategies are defined in mapping data
 - Mapping data is always used by all use cases
 - Join Fetch Queries are defined in code
 - Only executed in the use case that it is defined in
- Like Eager Joining, join fetch queries use SQL joins to pre-cache additional data
 - Extra data is not returned as part of the result set



Join Fetch Queries

- Queries can safely join multiple referenced objects
- Should not join more than one collection
 - Even for a single collection 'distinct' is needed
 - Multiple collections create a Cartesian product ⚡

```
Query query = session.createQuery("select distinct p "  
    + "from Person p left join fetch p.accounts");  
List<Person> people = query.list();
```

Fetch joins are outer joins even if you do not specify LEFT or OUTER

Hibernate:
select

```
    distinct person0_.id as id0_0_,  
    accounts1_.number as number1_1_,  
    person0_.firstname as firstname0_0_,  
    person0_.lastname as lastname0_0_,  
    accounts1_.balance as balance1_1_,  
    accounts1_.owner_id as owner3_1_1_,  
    accounts1_.owner_id as owner3_0__,  
    accounts1_.number as number0__  
from  
    Person person0_  
left outer join  
    Account accounts1_  
    on person0_.id=accounts1_.owner_id
```

```
Criteria criteria = session.createCriteria(Person.class)  
    .setFetchMode("accounts", FetchMode.JOIN)  
    .setResultTransformer(Criteria.DISTINCT_ROOT_ENTITY);  
List<Person> people = criteria.list();
```

Loads person objects and pre-cache the associated accounts using a single select



Hibernate Optimization:

STRATEGY: BATCH FETCHING



Batch Fetching

■ References N+1:

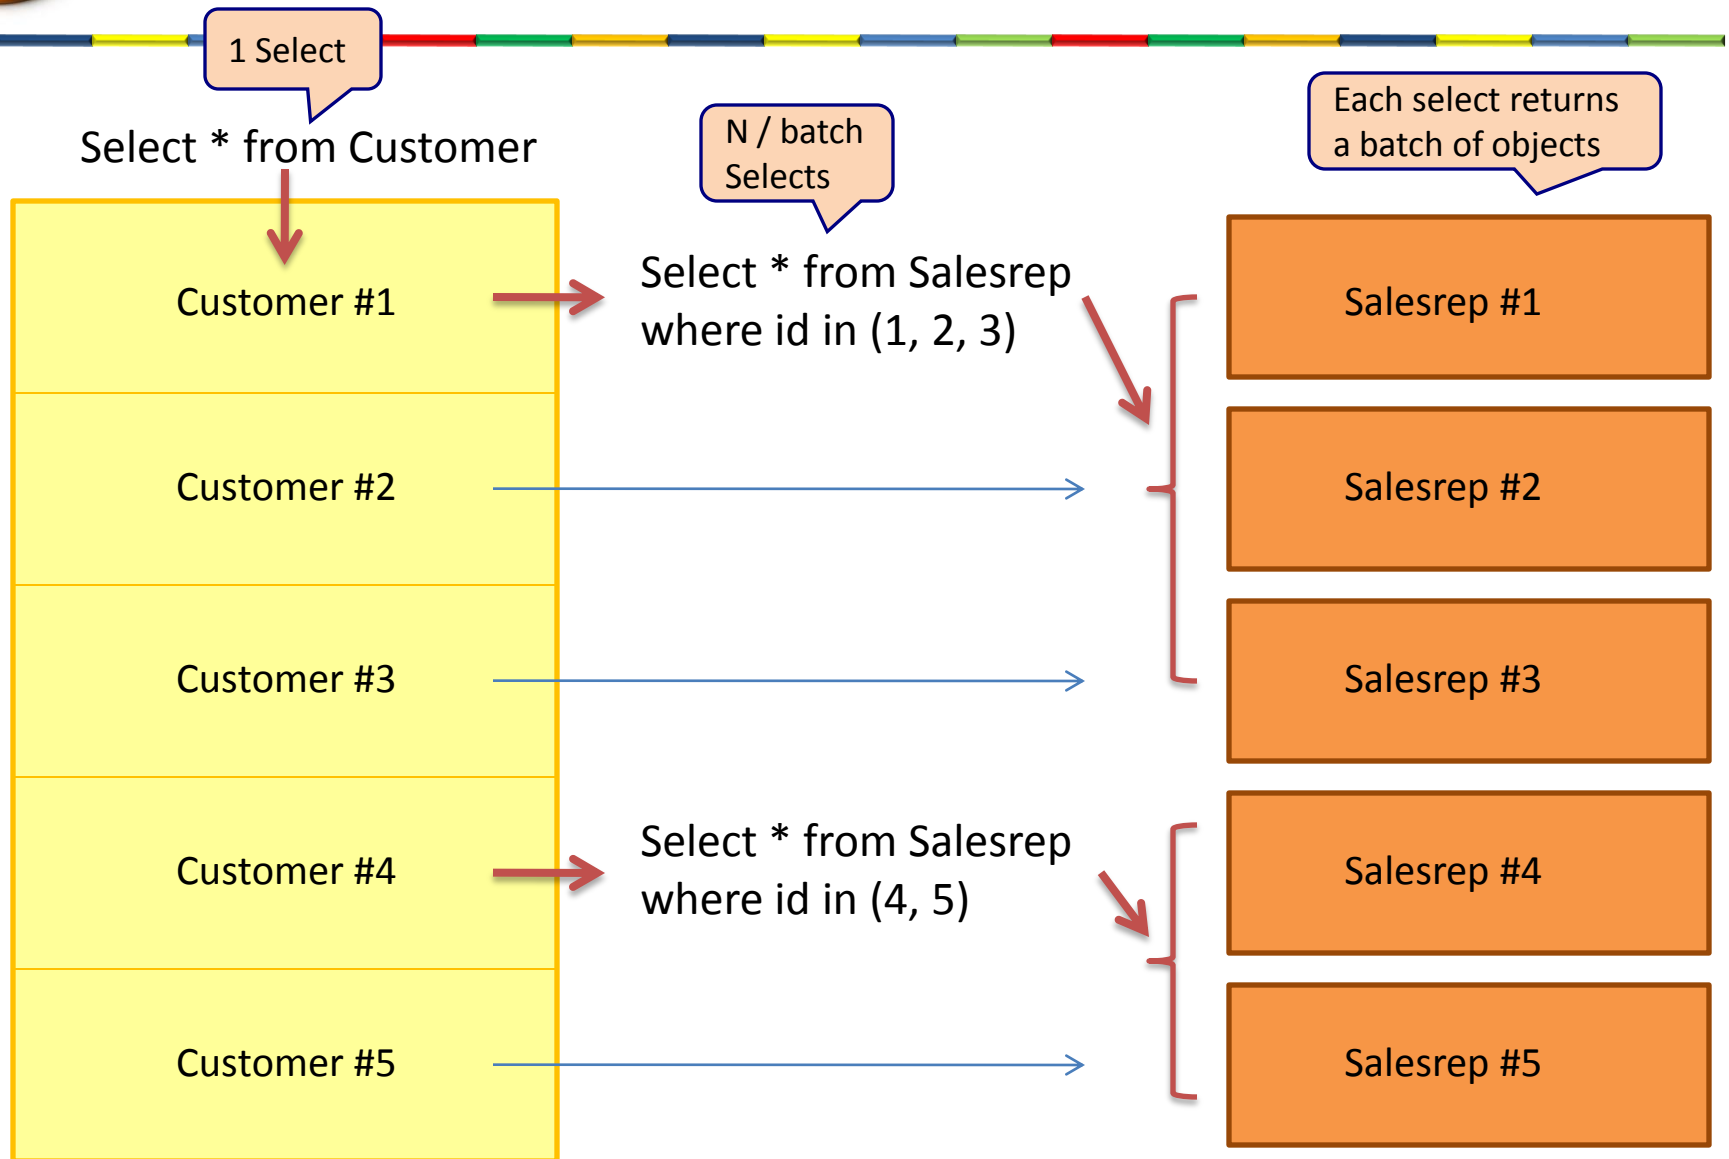
- We also saw how N+1 loads the salesrep for each customer in a separate select
- **Batch fetching helps:** by loading the salesrep for several customers in one go – loading a batch
 - When the first reference is needed

■ Collections N+1:

- We saw how N+1 loads the customer list for each salesrep using a separate select
- **Batch fetching helps:** by loading the customer lists for several salesreps in one go – loading a batch
 - When the first collection is needed



Batch Fetching References





Batch References

```
@Entity
public class Customer {
    @Id
    @GeneratedValue
    private int id;
    private String firstname;
    private String lastname;

    @ManyToOne(fetch = FetchType.LAZY)
    private SalesRep salesRep;

    ...
}
```

```
@Entity
@org.hibernate.annotations.BatchSize(size=3)
public class SalesRep {
    @Id
    @GeneratedValue
    private int id;
    private String name;

    @OneToMany(mappedBy="salesRep")
    private Set<Customer> customers = new HashSet<Customer>();

    ...
}
```

SalesRep will be loaded in batches of 3 or less, when possible

```
List<Customer> customers =
    session.createQuery("from Customer").list();
SalesRep salesrep = null;
for (Customer customer : customers) {
    salesrep = customer.getSalesRep();
    // do something with the salesrep
    salesrep.getName();
}
```

Batch fetching works because customers with un-retrieved salesrep have been loaded

Hibernate:

```
select
    customer0_.id as id0_,
    customer0_.firstname as firstname0_,
    customer0_.lastname as lastname0_,
    customer0_.salesRep_id as salesRep4_0_
from
    Customer customer0_

Hibernate:
select
    salesrep0_.id as id1_0_,
    salesrep0_.name as name1_0_
from
    SalesRep salesrep0_
where
    salesrep0_.id in (
        ?, ?, ?
    )
```



Batch References– XML

```
<hibernate-mapping >
  <class name="Customer">
    <id name="id">
      <generator class="native" />
    </id>
    <property name="firstname" />
    <property name="lastname" />
    <many-to-one name="salesRep" />
  </class>
</hibernate-mapping>
```

```
List<Customer> customers =
    session.createQuery("from Customer").list();
SalesRep salesrep = null;
for (Customer customer : customers) {
    salesrep = customer.getSalesRep();
    // do something with the salesrep
    salesrep.getName();
}
```

Batch fetching works when Hibernate knows about un-retrieved salereps

```
<hibernate-mapping package="how.always.batch.entity">
  <class name="SalesRep" batch-size="3">
    <id name="id">
      <generator class="native" />
    </id>
    <property name="name" />
    <set name="customers" inverse="true" cascade="persist">
      <key column="salesRep" />
      <one-to-many class="Customer" />
    </set>
  </class>
</hibernate-mapping>
```

batch-size attribute on the <class> tag

Hibernate:

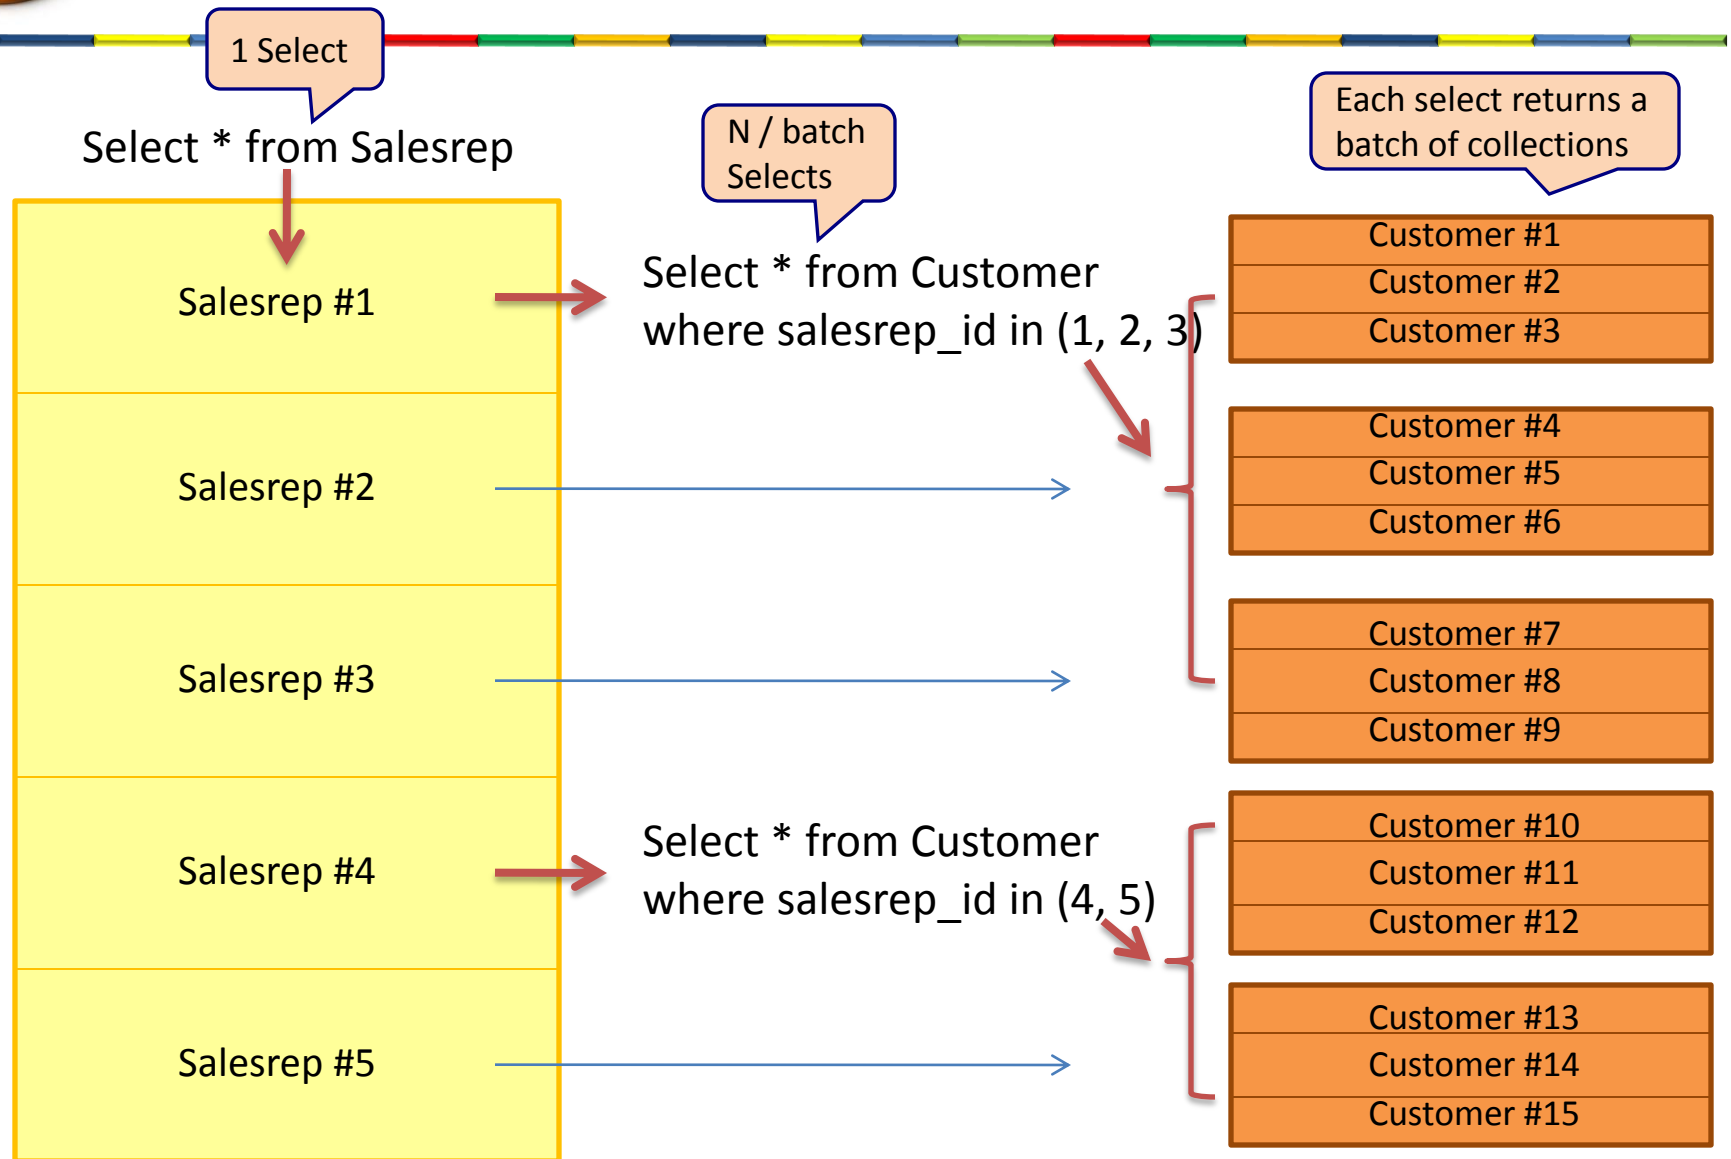
```
select
    customer0_.id as id0_,
    customer0_.firstname as firstname0_,
    customer0_.lastname as lastname0_,
    customer0_.salesRep_id as salesRep4_0_
from
    Customer customer0_
```

Hibernate:

```
select
    salesrep0_.id as id1_0_,
    salesrep0_.name as name1_0_
from
    SalesRep salesrep0_
where
    salesrep0_.id in (
        ?, ?, ?
    )
```




Batch Fetching Collections





Batch Fetching – Collections

```
@Entity
public class Customer {
    @Id
    @GeneratedValue
    private int id;
    private String firstname;
    private String lastname;

    @ManyToOne(fetch = FetchType.LAZY)
    private SalesRep salesRep;

    ...
}
```

```
@Entity
public class SalesRep {
    @Id
    @GeneratedValue
    private int id;
    private String name;

    @OneToMany(mappedBy="salesRep", cascade=CascadeType.PERSIST)
    @org.hibernate.annotations.BatchSize(size=3)
    private Set<Customer> customers = new HashSet<Customer>();

    ...
}
```

Try to load the customer collection for 3 salesreps when possible

```
List<SalesRep> salesreps =
    session.createQuery("from SalesRep").list();
Set<Customer> customers = null;
for (SalesRep s : salesreps) {
    customers = s.getCustomers();
    for (Customer c : customers) {
        // do something with the customer
    }
}
```

Batch fetching only works because
Hibernate knows of un-retrieved
customer collections



Hibernate:

```
select
    salesrep0_.id as id1_,
    salesrep0_.name as name1_
from
    SalesRep salesrep0_
```

Hibernate:

```
select
    customers0_.salesRep_id as salesRep4_1_,
    customers0_.id as id1_,
    customers0_.id as id0_0_,
    customers0_.firstname as firstname0_0_,
    customers0_.lastname as lastname0_0_,
    customers0_.salesRep_id as salesRep4_0_0_
from
    Customer customers0_
where
    customers0_.salesRep_id in (
        ?, ?, ?
    )
```



Batch Collections – XML

```
<hibernate-mapping >
  <class name="Customer">
    <id name="id">
      <generator class="native" />
    </id>
    <property name="firstname" />
    <property name="lastname" />
    <many-to-one name="salesRep" />
  </class>
</hibernate-mapping>
```

```
<hibernate-mapping package="how.always.batch.collection">
  <class name="SalesRep">
    <id name="id">
      <generator class="native" />
    </id>
    <property name="name" />
    <set name="customers" batch-size="3"
      inverse="true" cascade="persist">
      <key column="salesRep" />
      <one-to-many class="Customer" />
    </set>
  </class>
</hibernate-mapping>
```

XML uses the batch-size attribute on collection tags

Hibernate:

```
select
  salesrep0_.id as id1_,
  salesrep0_.name as name1_
from
  SalesRep salesrep0_
```

Hibernate:

```
select
  customers0_.salesRep_id as salesRep4_1_,
  customers0_.id as id1_,
  customers0_.id as id0_0_,
  customers0_.firstname as firstname0_0_,
  customers0_.lastname as lastname0_0_,
  customers0_.salesRep_id as salesRep4_0_0_
from
  Customer customers0_
where
  customers0_.salesRep_id in (
    ?, ?, ?
  )
```

```
List<SalesRep> salesreps =
  session.createQuery("from SalesRep").list();
Set<Customer> customers = null;
for (SalesRep s : salesreps) {
  customers = s.getCustomers();
  for (Customer c : customers) {
    // do something with the customer
  }
}
```

Batch fetching only works because
Hibernate knows of un-retrieved
customer collections





Batch Fetching

- Batch fetching is an easy and safe optimization
 - If un-needed data is retrieved it's never much
 - No joins are involved, no Cartesian Product
 - Can be specified for references and collections
- Typical batch sizes are between 3 and 15
- Batch fetching reduces the $N + 1$ problem to
→ $\text{Ceil}(N / \text{Batch Size}) + 1$



Hibernate Optimization:

STRATEGY: SUB SELECT

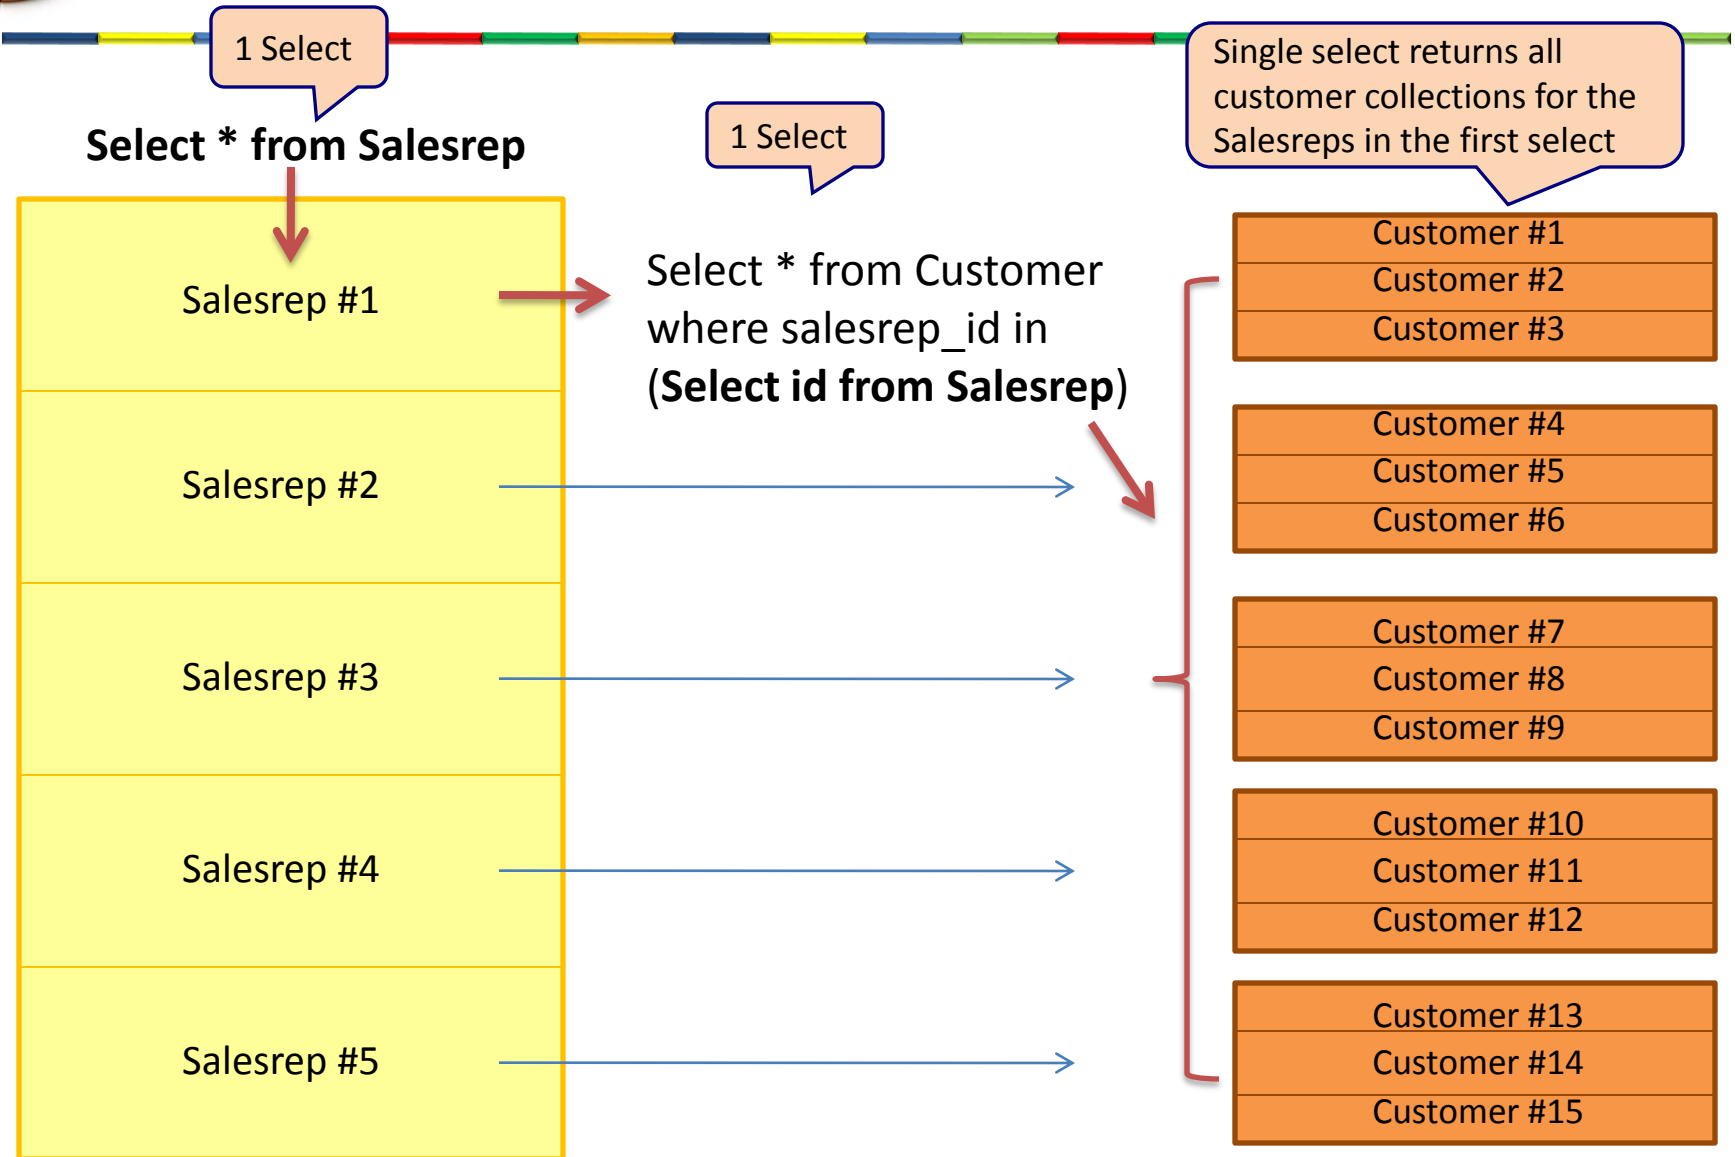


Sub Select

- The sub select strategy is a specialized form of the batch fetching strategy for collections
 - Instead of loading a batch of collections it loads **all related collections in one select**
 - Just like batch fetching it **doesn't retrieve** anything until the **first time a collection** is needed
- Sub select is not available for references



Batch Fetching Collections





Sub Select Collections

Keeps track of the query used to retrieve the salesreps

```
List<SalesRep> salesreps = session.createQuery(
    "from SalesRep where id < 1000").list();
Set<Customer> customers = null;
for (SalesRep s : salesreps) {
    customers = s.getCustomers();
    for (Customer c : customers) {
        // do something with the customer
    }
}
```

Sub-Select eager fetching only works for collections

```
@Entity
public class SalesRep {
    @Id
    @GeneratedValue
    private int id;
    private String name;

    @OneToMany(mappedBy="salesRep", cascade=CascadeType.PERSIST)
    @org.hibernate.annotations.Fetch(
        org.hibernate.annotations.FetchMode.SUBSELECT
    )
    private Set<Customer> customers = new HashSet<Customer>();
}
```

FetchMode.SUBSELECT

Hibernate:

```
select
    salesrep0_.id as id1_,
    salesrep0_.name as name1_
from
    SalesRep salesrep0_
where
    salesrep0_.id<1000
```

Hibernate:

```
select
    customers0_.salesRep_id as salesRep4_1_,
    customers0_.id as id1_,
    customers0_.id as id0_0_,
    customers0_.firstname as firstname0_0_,
    customers0_.lastname as lastname0_0_,
    customers0_.salesRep_id as salesRep4_0_0_
from
    Customer customers0_
where
    customers0_.salesRep_id in (
        select
            salesrep0_.id
        from
            SalesRep salesrep0_
        where
            salesrep0_.id<1000
    )
```

Re-uses that query as a sub select to get the customer collections for those salesreps




Sub Select – XML

Keeps track of the query used to retrieve the salesreps

```
List<SalesRep> salesreps = session.createQuery(
    "from SalesRep where id < 1000").list();
Set<Customer> customers = null;
for (SalesRep s : salesreps) {
    customers = s.getCustomers();
    for (Customer c : customers) {
        // do something with the customer
    }
}
```

Sub-Select eager fetching only works for collections



Hibernate:

```
select
    salesrep0_.id as id1_,
    salesrep0_.name as name1_
from
    SalesRep salesrep0_
where
    salesrep0_.id<1000
```

Hibernate:

```
select
    customers0_.salesRep_id as salesRep4_1_,
    customers0_.id as id1_,
    customers0_.id as id0_0_,
    customers0_.firstname as firstname0_0_,
    customers0_.lastname as lastname0_0_,
    customers0_.salesRep_id as salesRep4_0_0_
from
    Customer customers0_
where
    customers0_.salesRep_id in (
        select
            salesrep0_.id
        from
            SalesRep salesrep0_
        where
            salesrep0_.id<1000
    )
```

Re-uses that query as a sub select to get the customer collections for those salesreps

Fetch="subselect"

```
<hibernate-mapping package="how.always.subselect">
  <class name="SalesRep">
    <id name="id">
      <generator class="native" />
    </id>
    <property name="name" />
    <set name="customers" fetch="subselect"
        inverse="true" cascade="persist">
      <key column="salesRep" />
      <one-to-many class="Customer" />
    </set>
  </class>
</hibernate-mapping>
```



Sub Select

- The **Sub Select** strategy solves the $N + 1$ problem by **turning it into a $1 + 1$**
 - Only available for collections, not references
 - May retrieve too much data if you did not actually need to work with the collections
 - Like batch fetching, no joins, no Cartesian Product
- Internally Sub Select keeps track of the query used to retrieve the original objects





Hibernate Optimization:

2ND LEVEL CACHING



2nd Level Caching

- By default Hibernate only uses Session Caches
 - Objects are cached for the duration of the session
- You can enable a second level cache
 - Lasts for the duration of the SessionFactory
 - Shared by all sessions





Caching and Optimization

- 2nd Level caching should never be used as an alternative to fetch optimizations
 - Can not solve problems, can attempt to hide them
 - Should be used to help scale the application
- Caching is a large and complex field
 - We will cover Hibernates basic caching features
 - Improper configuration can create problems that are difficult to debug





What to Cache



- Hibernate can cache entity objects and collections (collections of entity IDs)
 - But not all of them will benefit from being cached
- Good candidates for caching :
 - Do not change, or change rarely
 - Are modified only by your application
 - Are non-critical to the application
- Typical examples include reference data
 - Such as customer categories, or statuses



Caching Strategies

- Four different caching strategies:
 - **Read Only**: very fast caching strategy, but can only be used for data that never changes
 - **Non Strict Read-Write**: data may be stale for a while, but it does get refreshed at timeout
 - **Read-Write**: prevents stale data, but at a cost. Use for read-mostly data in a non-clustered setup
 - **Transactional**: Can prevent stale data in a clustered environment. Can be used for read-mostly data

Stricter and therefore Slower






Cache Providers

- The following open source cache providers are bundled with Hibernate
- Only a single cache provider per SessionFactory

Provider	Read Only	Non Strict Read Write	Read Write	Transactional
EHCache	✓	✓	✓	
OSCache	✓	✓	✓	
SwarmCache	✓	✓		
JBoss Cache 1.x	✓			✓
JBoss Cache 2.x	✓			✓




Caching Example – Entities



- **Category entities – Read Only**
 - Typical reference data, categories are never updated
- **SalesRep entities – Non Strict Read Write**
 - Not many Salesreps, always needed when editing customers
 - SalesReps seldom change, stale SalesRep records are fine
- **Customer entities – Not Cached**
 - Too many customers, customer are updated frequently



Caching Example – Collections



- **Customer Collection for each category – Read Write**
 - Often used, try to avoid stale data as much as possible
- **Customer Collection for each SalesRep – Not Cached**
 - Not used frequently enough to warrant caching



Category

```
@Entity
@org.hibernate.annotations.Entity (mutable=false)
@org.hibernate.annotations.Cache (usage=
    CacheConcurrencyStrategy.READ_ONLY
)
public class Category {
    @Id
    private String abbreviation;
    private String description;

    @OneToMany (mappedBy="category")
    @org.hibernate.annotations.Cache (usage=
        CacheConcurrencyStrategy.READ_WRITE
    )
    private Set<Customer> customers = new HashSet<Customer>();

    ...
}
```

Mutable=false indicates to Hibernate that Categories can never change

Specify read only caching for Category Entities

Specify read write caching for the collection of customers each category has

```
<hibernate-mapping package="cacheDemo">
  <class name="Category" mutable="false">
    <cache usage="read-only" />
    <id name="abbreviation" />
    <property name="description" />
    <set name="customers" inverse="true" cascade="persist">
      <cache usage="read-write" />
      <key column="salesRep" />
      <one-to-many class="Customer" />
    </set>
  </class>
</hibernate-mapping>
```

Mutable=false inside <class> tag

XML uses <cache> tag inside <class> to specify category entity caching

<cache> tag inside <set> for the customers collection



SalesRep

```
@Entity
@org.hibernate.annotations.Cache (usage=
    CacheConcurrencyStrategy.NONSTRICT_READ_WRITE
)
public class SalesRep {
    @Id
    @GeneratedValue
    private int id;
    private String name;

    @OneToMany (mappedBy="salesRep", cascade=CascadeType.PERSIST)
    private Set<Customer> customers = new HashSet<Customer>();

    ...
}
```

Specify non strict read write caching for SalesRep Entities

```
<hibernate-mapping package="cacheDemo">
  <class name="SalesRep">
    <cache usage="nonstrict-read-write" />
    <id name="id">
      <generator class="native" />
    </id>
    <property name="name" />
    <set name="customers" inverse="true" cascade="persist">
      <key column="salesRep" />
      <one-to-many class="Customer" />
    </set>
  </class>
</hibernate-mapping>
```

XML uses the <cache> tag inside the <class> tag



Enabling Caching (EHCache)

```
<hibernate-configuration>
  <session-factory>
    <!-- HSQL DB running on localhost -->
    <property name="connection.url">jdbc:hsqldb:hsq1://localhost/trainingdb</property>
    <property name="connection.driver_class">org.hsqldb.jdbcDriver</property>
    <property name="connection.username">sa</property>
    <property name="connection.password"></property>
    <property name="dialect">org.hibernate.dialect.HSQLDialect</property>

    <!-- Enable Second Level Cache -->
    <property name="cache.provider_class">org.hibernate.cache.EhCacheProvider</property>

    <!-- Enable Statistics -->
    <property name="generate_statistics">true</property>

    <!-- Hibernate XML mapping files - Cache -->
    <mapping resource="cacheDemo/Customr.hbm.xml" />
    <mapping resource="cacheDemo/SalesRep.hbm.xml" />
    <mapping resource="cacheDemo/Category.hbm.xml" />
  </session-factory>
</hibernate-configuration>
```

Enable 2nd level caching by specifying a caching provider

Optionally enable statistics



Configuring EHCACHE – Cache Eviction

```
<ehcache>
  <diskStore path="java.io.tmpdir"/>
  <defaultCache
    maxElementsInMemory="10000"
    eternal="false"
    timeToIdleSeconds="120"
    timeToLiveSeconds="120"
    overflowToDisk="true" />
```

EHCache General configuration

```
<cache name="cacheDemo.Category"
  maxElementsInMemory="50"
  eternal="true"
  timeToIdleSeconds="0"
  timeToLiveSeconds="0"
  overflowToDisk="false" />
```

Sets up a cache region for category entities

```
<cache name="cacheDemo.Category.customers"
  maxElementsInMemory="50"
  eternal="false"
  timeToIdleSeconds="3600"
  timeToLiveSeconds="7200"
  overflowToDisk="false" />
```

Sets up a cache region for the customer collections inside the category entities

```
<cache name="cacheDemo.SalesRep"
  maxElementsInMemory="500"
  eternal="false"
  timeToIdleSeconds="1800"
  timeToLiveSeconds="10800"
  overflowToDisk="false" />
```

cache region for the SalesRep entities

```
</ehcache>
```



Hibernate Statistics

Stats object also holds
general statistics for many
Hibernate subsystems

```
Statistics stats = sessionFactory.getStatistics();  
long hits    = stats.getSecondLevelCacheHitCount();  
long misses  = stats.getSecondLevelCacheMissCount();  
long puts    = stats.getSecondLevelCachePutCount();  
System.out.printf("\nGeneral 2nd Level Cache Stats\n");  
System.out.printf("Hit: %d Miss: %d Put: %d\n", hits, misses, puts);
```

General 2nd level cache statistics

```
SecondLevelCacheStatistics salesRepStats =  
    stats.getSecondLevelCacheStatistics("cacheDemo.SalesRep");  
long srCurrent = salesRepStats.getElementCountInMemory();  
long srMemsize = salesRepStats.getSizeInMemory();  
long srHits    = salesRepStats.getHitCount();  
long srMisses  = salesRepStats.getMissCount();  
long srPuts    = salesRepStats.getPutCount();  
System.out.printf("\nSalesRep Cache Region - Size: %d Holds: %d\n", srMemsize, srCurrent);  
System.out.printf("Hit: %d Miss: %d Put: %d\n", srHits, srMisses, srPuts);
```

cache statistics for a
specific cache region

```
Statistics stats = sessionFactory.getStatistics();  
Stats.clear();  
stats.setStatisticsEnabled(true);  
  
...  
  
stats.setStatisticsEnabled(false);
```

Statistics can also be enabled
or disabled programmatically
allowing you to do more
targeted measurements



Hibernate Optimization:

WRAPPING UP



Analyze SQL

- Before changing any fetching strategies
 - Analyze the SQL Hibernate uses for all use cases
 - Look for things that can actually cause problems
 - Don't over optimize, only update real problem areas
- Then after each change check the SQL again

```
<hibernate-configuration>
  <session-factory>

    <property name="show_sql">true</property>
    <property name="format_sql">true</property>
    <property name="use_sql_comments">true</property>

    ...

  </session-factory>
</hibernate-configuration>
```

The following three property can be used to check Hibernate's SQL



Optimization Steps

1. Update Fetching Strategies by:
 - a) Analyzing SQL used by Hibernate
 - b) Updating a Fetching Strategy
 - c) Checking SQL to ensure effectiveness
2. Enable Second Level Caching
 - Fine-tuning the cache settings in production



Active Learning

- Describe the difference between batch fetching and sub select optimization.
- Why doesn't second level caching fix bad fetching strategies?



Module Summary

- Data Access Optimization changes when and how Hibernate retrieves data
- Hibernate mostly defaults to lazy loading
 - Lazy loading can lead to too many small selects
 - Incorrect eager loading can lead to slow queries
- 2nd level caching should not be used as an alternative to fetch optimizations
 - Caching can help boost performance under load
 - Incorrectly configured cache can create problems